


H3 PROMPT COMPOSER

Illustrated User Guide

V5.11.1

What this is A standalone local HTML tool that turns plain-language creative choices into structured MiniMax H3 prompts. You choose references, shots, camera behavior, dialogue, timing, and continuity; the composer handles the prompt structure and checks for common conflicts.

This guide assumes you can already run MiniMax H3 in ComfyUI, but it does not assume you know H3 prompt syntax. You do not need to memorize <Subject>, <Picture>, <Video>, or <Audio> rules to use the composer.



The four areas you will use most: setup, Generations, Shot Builder, and the generated prompt/check panel.

1. Quick Start

The shortest path from a blank page to a usable H3 prompt.

- 1 **Choose a mode** - Use Ref2VA for most reference-heavy H3 work. Use T2VA/I2VA/FL2VA/L2VA for the base H3 modes.
- 2 **Set duration and style** - Enter the target seconds and a short visual-style description.
- 3 **Add references** - For Ref2VA, add reusable characters/subjects, an environment, and any voice/audio/video references you actually need.
- 4 **Build the Generation** - A Generation is one H3 output. Add one or more Shots only when the camera actually cuts.
- 5 **Describe each Shot** - Set the environment, opening placement/action, optional Timed Action Beats, camera, dialogue, and sound.
- 6 **Read Prompt Check** - Fix red errors. Review amber warnings. Blue information messages are guidance, not failures.
- 7 **Copy the prompt** - Paste the generated prompt into the matching H3 conditioning node in ComfyUI.

Best first move Open Load example and choose Full reference scene or Simple text-to-video shot. Keep that example open in a second browser window while you build your own project.


The composer is not a media loader

The HTML does not load your ComfyUI images, videos, or audio. It describes how those already-connected ComfyUI inputs should be interpreted. If Picture 3 is selected in the composer, Picture input 3 in your H3 node should contain the matching image.

2. The Mental Model

Most confusion disappears once you separate Generation, Shot, and Timed Action Beat.

Level	What it means	Create a new one when...
Generation	One complete H3 output / prompt. References are shared across Generations.	You are making the next clip/output in the sequence.
Shot	One continuous camera view inside that Generation.	There is an actual cut, transition, or new camera view.
Timed Action Beat	A timed action window inside one continuous Shot.	The camera does not cut, but Action A must happen before/for a different amount of time than Action B.



A Generation contains one or more Shots. Shot 1 opens at 0:00; later Shots get cut timestamps.

Do not create fake Shots for pacing If a 10-second camera view stays continuous, keep one Shot. Use Timed Action Beats to control what happens during 0-5 seconds versus 5-10 seconds.

3. Choose the Right Mode

Pick the H3 task first. The rest of the interface adapts around it.

Mode	Use it for	Image relationship
T2VA	Text-only generation	No image alignment instruction.
I2VA	Animate from a starting image	Picture 1 is the first frame.
FL2VA	Travel from a first image to a last image	Picture 1 opens; Picture 2 is the final frame.
L2VA	Generate toward a known ending image	Picture 1 is the final frame.
Ref2VA	Characters, environments, voices, videos, editing, multi-reference generation	Use Subjects plus Picture/Video/Audio references.

Most reference work = Ref2VA If you are using character sheets, multiple environment images, voice references, performance videos, or edit workflows, Ref2VA is usually the right starting point.

Basic vs. Advanced

Basic view exposes the controls most people need. Advanced view adds lower-level reference/task options and specialized camera/reference controls. Start in Basic; switch to Advanced only when you know what you need to change.

4. Ref2VA References

Give each reference one clear job. The composer turns those jobs into H3 labels automatically.

A Subject is the reusable thing you want H3 to keep track of: a person, creature, vehicle, prop, or environment. The source image/video that defines it is attached inside that Subject. A standalone Picture is reserved for an actual frame/composition anchor, storyboard frame, or special scale reference.

Environment: multiple source images can describe the same stable location. Voice: bind one audio input to one Subject; Auto sends it only when that Subject speaks.

Reference numbering Picture and Video numbers correspond to the physical H3 input slots. Audio labels are generated from the active audio references in H3 order, so bypassed voices do not consume an <Audio N> label. The Identity and Voice Map shows what will actually be sent.

5. Environments, Voice Mapping, and Scale

Three areas that strongly affect continuity.

Environments

Attach every useful view of the same location to a single Environment Subject. The prompt treats those images as complementary views of one stable environment and asks H3 to preserve architecture, layout, furnishings, materials, colors, and lighting. Shot framing still comes from the Shot Builder, not from the environment source image.

Voice mapping

SUBJECT	SPEAKER	VOICE
<Subject 1> Maya	(S1)	<Audio 1>
<Subject 2> Eli	(S2)	<Audio 2>

H3 LABEL	SOURCE	PURPOSE
<Audio 1>	Audio input 1	voice-timbre reference for <Subject 1>
<Audio 2>	Audio input 2	voice-timbre reference for <Subject 2>

Audio labels restart for each Generation. H3 numbers synchronized Video soundtracks first, followed by separate Audio inputs in socket order.

This table is for you. The generated prompt does not add a separate global voice-mapping sentence; it applies each Audio reference locally

The map is for you: it shows which Subject speaks as which speaker ID and which active H3 Audio label is used.

Bind a voice once at the reference level. In the Dialogue card, selecting that Subject automatically applies the matching voice reference. This avoids manually typing <Audio 1>, <Audio 2>, and so on into every line.

Height and scale

Maya
 <Picture 5> height reference
 auto definition

What does this image control?

Visual height / scale — global attribute source

This Picture is a global attribute source, not a shot reference. Only height, body scale, proportions, and floor-plane information are transferred. It is never assigned to a Shot or repeated in the shot timeline, and the prompt explicitly tells H3 not to reproduce or transition from its composition. If the model still imitates the image, switch to Text-only height guidance and remove this image from the connected inputs.

Short name

Maya and Eli height reference

Use a deliberately utilitarian scale diagram when possible: neutral poses, shared floor baseline, and minimal cinematic composition. This Picture is global size guidance only; it never controls a target frame or blocking.

Connected Picture used for this scale reference

Picture 5

Choose the actual ComfyUI Picture input containing the scale diagram. You can change this at any time; creation order does not determine the reference.

Character shown on the left

<Subject 1> — Maya

Character shown on the right (optional)

<Subject 2> — Eli

Additional scale notes (optional)

Preserve only the visible height difference; ignore the diagram's pose, spacing, and camera angle.

Subject <Subject 3> apartment kitchen auto definition

Subject type

Environment / location

Name

apartment kitchen

Reference contribution

Environment identity and layout

This tells H3 what to take from the Subject's attached images and videos; shot framing still comes from the Shot and storyboard controls.

Environment source images — select one or more views of the same location 2 views selected

Picture 1 Picture 2 Picture 3 Picture 4 Picture 5

Select every image that shows a complementary view of this same environment. All selected images define one Environment Subject; they are not treated as separate locations.

Environment source videos

Set the Videos count above to attach a source video.

Attach image views or source footage of the same place. Environment references establish architecture and spatial continuity without forcing their camera framing into new shots.

Environment details to preserve

Warm apartment kitchen with white cabinets, a wooden island, two pendant lights, a window over the sink, an open doorway to the hall, and soft

Visual scale reference: use only when text-only height guidance is not strong enough.

Environment references define the place; they do not define the final camera framing.

Prefer the simplest authority Start with text-only height relationships when possible. A scale image is stronger, but any extra image gives the model more visual information to reconcile.

6. Build a Shot

Use the normal action field for the opening state; use Timed Action Beats only when pacing matters.

The normal Subject placement and action field is best for what is true at the start of the Shot: where the subject is, what they are holding, where they are looking, or an untimed action. If every action is already fully described by Timed Action Beats, the normal action field can be left empty.

Shot 1 opens — no timestamp – ↑ ↓ Duplicate + On-screen text + Dialogue

Subject placement and action before dialogue

A ceramic artist sits at a wooden wheel in a sunlit studio with both hands resting on the centered lump of wet clay.

Use this for the opening placement/state and any untimed action. For multiple actions that need specific pacing within the same continuous Shot, use Timed Action Beats below. Put reactions after a spoken line inside that Dialogue card's "Action after this line" field.

Timed Action Beats — optional 2 beats

Use this when several actions must occupy specific parts of one continuous Shot. Times are relative to this Shot and do not create cuts. Use the normal action field above for the opening placement or state. If the beats fill the whole Shot and the camera moves at the same time, set Camera Motion timing to Throughout the shot (or another simultaneous timing) rather than After opening action.

Planned Shot duration: 8.00 seconds

BEAT 1 0.00–4.00 sec ×

Start — seconds into this Shot

End — seconds into this Shot

4

Action during this time window

She presses both thumbs into the spinning clay and steadily raises the walls of the vessel.

BEAT 2 4.00–8.00 sec ×

Start — seconds into this Shot

End — seconds into this Shot

8

Action during this time window

She eases the rim outward, then pauses to inspect the finished shape.

+ Add beat Distribute evenly

One continuous Shot: opening state above, then two timed actions below. The times are relative to this Shot and do not create cuts.

Example: 10-second Shot Beat 1 = 0.00-5.00 seconds; Beat 2 = 5.00-10.00 seconds. Use Distribute evenly when you simply want equal time for each beat.

7. Camera Builder

Think in endpoints: where the camera starts, where it ends, and how it gets there.

Camera Builder Start: Wide / full shot · Front three-quarter — viewing more of subject's left side · End: Medium ...

1 START FRAME

Framing / shot size

Wide / full shot

Subject or Environment being framed

— automatic / primary target —

Second subject in frame — optional

— none / single-subject framing —

Choose a second character here to make the selected shot size a two-person framing, such as a medium two-shot.

Horizontal viewpoint

Front three-quarter — viewing n

Camera height and vertical angle

Slightly high, looking gently down

Composition — optional

Automatic / unspecified

2 END FRAME

Every blank End control inherits Start if a distinct endpoint is defined. If all End controls remain on "Same as Start," Step 3 can still use a loose movement range without forcing an exact ending frame.

Ending framing / shot size

Medium shot

Subject or Environment being framed

Same as Start — opening frame

Second subject in frame — optional

Same as Start — none / single-

Keep the same second subject, switch to another, or choose None to resolve from two people to a single subject.

Horizontal viewpoint

Same as Start — Front three-qu

Camera height and vertical angle

Same as Start — Slightly high, l

Composition — optional

Same as Start — thirds

3 CAMERA MOTION

Endpoint analysis: tighter ending framing. Choose how the camera physically gets from Frame 1 to Frame 2.

Movement

Push In

Recommended: Push In naturally connects the wider Start Frame to the tighter End Frame.

Movement target

Automatic — follow Start Fi

Leave this on Automatic for normal subject-focused moves. Choose a different reference only when the camera is intentionally anchored to something other than the Start Frame target.

Custom target instead

the ceramic artist

Speed

Slow

Motion timing

Throughout the shot

Controls when the movement happens relative to the shot action and dialogue; it changes prompt order instead of adding redundant timing instructions.

Additional movement instruction — optional

reveal <Subject 2> in the doorv

Describe only something not already established by the End Frame.

► Stability, lens, and focus

CAMERA INSTRUCTION PREVIEW

A wide shot shows the subject in full with the immediate environment, with the camera at a front three-quarter angle favoring the subject's left side. The camera is slightly above eye level, looking gently downward. Throughout the shot, the camera slowly pushes in toward the ceramic artist, tightening to a medium shot, with smooth stabilized movement.


Camera Builder: 1) Start Frame, 2) End Frame, 3) Camera Motion. The preview at the bottom shows the exact camera instruction being generated.

- Start Frame: define the opening size, target, viewpoint, height, and composition.
- End Frame: leave Same as Start unless the camera must finish differently.
- Camera Motion: choose a physically compatible move and when it happens. If Timed Beats fill the whole Shot, use Throughout the shot when the camera moves during them.


Avoid duplicate camera language Let Camera Builder own framing/motion whenever possible; use freeform action text only for details it cannot express.

8. Dialogue and Audio

The composer keeps speaker identity, voice reference, spoken words, and post-line action in timeline order.



Dialogue card: select the speaker, enter spoken words, and optionally add an action after the line.



Voice reference: bind the audio once to the Subject and choose Auto/Always/Off for the Generation.

- Spoken words only: enter the exact line without quotation marks or extra direction.
- Delivery: use the dialogue settings for tone/performance direction rather than embedding stage directions in the spoken words.
- Action after this line: put a reaction or action here when it must happen after the dialogue, so the compiler keeps the chronology correct.
- Soundscape: ambience and physical sounds for the scene. Fresh Generations inherit the base soundscape; Generation-specific extra sounds do not automatically carry forward.
- Non-diegetic music: audience-only score. Describe instrumentation, tempo, rhythm, and dynamics rather than only abstract mood words.

Why a voice may disappear from the prompt In Auto mode, a voice with no dialogue in the active Generation is bypassed. That is intentional: it does not consume an H3 <Audio N> label.

9. Generations and Continuity

Use a fresh Generation for a new output; use Duplicate when you want to branch from the same setup.

Control	What it does
+ New generation	Creates a fresh output while carrying relevant continuity such as the current appearance state, base soundscape, and music.
Duplicate	Copies the current Generation, including its Shots and settings, so you can make a variation without rebuilding it.
Appearance State	Stores a reusable look such as Jacket On, Jacket Off, Wet, Injured, etc.
Appearance change during this Shot	Transitions from one saved appearance state to another on camera; the ending state carries forward.

Historical safety When you change a shared Appearance State later, the composer can fork a variation for the current Generation instead of silently changing earlier prompts. Use the explicit edit-everywhere option only when you truly want a global change.

When to make another Shot instead

Create another Shot only when the camera actually cuts or transitions to a new view. If the same camera continues and only the subject behavior changes, stay in the same Shot and use Timed Action Beats.

10. Built-in Video Workflows

Use the guided starting points when you are editing or transferring from existing footage.

Workflow	What the composer preserves / transfers
Insert a subject into existing footage	Preserves the source camera/plate and helps describe placement, scale, occlusion, grounding, focus, motion blur, and source audio.
Replace background + relight	Preserves the performer motion while replacing the background and integrating environment lighting onto the performer.
Transfer performance + camera	Uses a filmed performance/camera reference to guide facial acting, mouth performance, gesture/body motion, exact performed dialogue, and/or camera movement.


Source-video camera ownership In direct edit workflows, the source video owns the camera. Camera Builder settings are intentionally suppressed so they do not fight the existing footage.

Load examples before building from scratch


The Load example menu contains full reference, insert-subject, background-relight, performance-transfer, and simple T2VA examples. These are useful as working templates and as a quick way to learn how the fields relate to the final prompt.

11. Prompt Check

Treat it like a preflight panel: red means fix, amber means review, blue means information.



Clean state: information messages plus a green “no blocking issues” result.




Example warning: the timed actions fill the Shot, but the camera is scheduled to move afterward.

Status	Meaning	What to do
Error	The prompt/setup is invalid or internally contradictory.	Fix before generating.
Warning	The prompt can still build, but the requested timing/reference/camera relationship deserves attention.	Review it; change it when the warning describes something you did not intend.
Info	Helpful context such as frame-grid snapping, word-count guidance, or an automatically bypassed voice.	Usually no action needed.

Word-count messages are guidance, not a target. Prefer concrete visible and audible information over filler.

12. Copy, Save, and Share

The right panel is the final product. The left/middle UI is just a safer way to build it.



The Prompt tab shows the exact text to paste into H3. Copy prompt copies the active Generation; Copy all generations copies every Generation in sequence.

- Copy prompt: copies the active Generation prompt.
- Copy all generations: copies every Generation prompt together, useful for planning a multi-clip sequence.
- Save .json: exports the editable composer project, not just the prompt text.
- Open: loads a previously saved composer JSON.
- Auto-saved locally: your current project is stored in browser local storage on that machine/browser profile.
- Restore previous: restores the safety backup created before destructive resets/preset changes.

To share a specific editable setup, export its .json file along with the HTML. Browser local storage does not travel with the file.

13. Troubleshooting and Cheat Sheet

When something looks wrong, simplify one variable at a time.

Symptom	First things to check
Wrong voice	Check Identity and Voice Map; confirm the correct ComfyUI Audio input is connected; confirm the intended Subject actually speaks in this Generation.
Character appears unexpectedly	Check whether the Subject is mentioned in any Shot, Timed Beat, dialogue, or transition. Remove inactive references from the current Generation when possible.
Characters merge / swap	Keep character references visually distinct; avoid redundant or conflicting identity sources; keep the anti-merge option enabled for multi-character shots.
Environment changes	Confirm all attached environment images really show the same location. Reduce conflicting references if needed. If using acceleration nodes, compare the same seed with them disabled or at conservative settings.
Second action gets rushed	Use Timed Action Beats instead of one long untimed action paragraph.
Camera warning with Timed Beats	If beats occupy the full Shot and the camera should move during them, choose Throughout the shot or During opening action.
Prompt feels too long	Remove repeated statements first. Define a reference once, describe the actual timeline, and avoid explaining the compiler to H3.

Five rules worth remembering

- Define each reusable thing once where it belongs.
- Give every reference one clear role.
- Use Shots for cuts; use Timed Action Beats for pacing inside one shot.
- Let Camera Builder own camera language whenever possible.
- Run Prompt Check before you copy the prompt.

If you are unsure Load the closest built-in example, replace its references/actions with yours, and change one section at a time. That is usually faster and safer than building a complicated Ref2VA prompt from a blank page.