H3 Prompt Composer

V5.11.1 User Guide

AUDITED SHARING BUILD

Contents

1. Quick Start

For most projects, this is the whole workflow:

1. **Open the HTML file** in a modern browser.
2. **Choose the H3 mode** you are using: T2VA, I2VA, FL2VA, L2VA, or Ref2VA.
3. **Set the target duration.**
4. In Ref2VA, **match the Images / Videos / Audios counts to the reference inputs connected in ComfyUI.**
5. **Define your Subjects and references -** characters, environments, voices, videos, scale references, and any frame anchors you need.
6. Build each **Generation** out of one or more **Shots**. Use the Camera Builder, dialogue cards, and Timed Action Beats as needed.
7. Read **Prompt Check**. Fix errors, review warnings, and treat information messages as guidance rather than failures.
8. Use **Copy Prompt** and paste the result into the appropriate MiniMax H3 prompt field in ComfyUI.

> **TIP:** Start simple. Only add a reference or control when it has a clear job. H3 generally benefits from a prompt where every instruction has one purpose.

---

2. The Three Timeline Levels

Understanding these three terms makes the rest of the composer much easier.

Generation

A **Generation** is one complete video generation request. It has its own Shots, summary, soundscape, music, active audio references, and appearance state.

Use **New Generation** when you want to create the next clip while carrying useful continuity forward. Use **Duplicate Generation** when you want another version of the same setup.

Shot

A **Shot** is a continuous camera setup. Adding another Shot means the video cuts or transitions to a new view.

Examples:

- Shot 1: Medium shot of Dani at the table.
- Shot 2 at 00:05.200: Cut to Ashley in the doorway.

Do **not** create a second Shot just because a character performs a second action while the camera remains continuous.

Timed Action Beat

A **Timed Action Beat** is an action window *inside the same Shot*.

For a continuous 10-second Shot:

```
0.00 - 5.00 sec   Dani folds the clothes and places them in the suitcase.
5.00 - 10.00 sec  Dani closes the suitcase, stands, and carries it to the door.
```

That remains one `[Shot 1]`. No cut is created.

---

3. Choosing a Mode

Mode	Use it when...	Main idea
T2VA	You are generating from text without an image keyframe.	Build the complete audiovisual timeline from text.
I2VA	Picture 1 is the exact first frame.	Begin at Picture 1 and develop forward.
FL2VA	Picture 1 is the first frame and Picture 2 is the last frame.	Describe a continuous path between the two frames.
L2VA	Picture 1 is the required final frame.	Build toward the final image.
Ref2VA	You are using character, environment, video, or audio references.	Define references once, then describe how they are used in the target video.

The composer automatically changes the output schema for the selected mode.

Base modes

T2VA, I2VA, FL2VA, and L2VA output the three core fields:

```
integrated_multimodal_description:
overall_soundscape:
non_diegetic_music:
```

I2VA, FL2VA, and L2VA also receive the correct first/last-frame alignment instruction.

Ref2VA

Full-reference mode outputs six sections:

```
subject_definitions:
summary:
retention_analysis:
detailed_description:
overall_soundscape:
non_diegetic_music:
```

You do not need to manually format those sections. The composer builds them from your choices.

---

4. Setting Up Ref2VA References

At the top of a Ref2VA project, set the connected **Images, Videos, and Audios** counts to match your ComfyUI workflow.

This tells the composer which physical reference inputs actually exist.

A note about numbering

ComfyUI may display physical inputs with zero-based socket names such as `ref_image_0` or `ref_audio_0`. H3 prompt references use one-based labels such as `<Picture 1>` and `<Audio 1>`.

For Pictures and Videos, the composer keeps the prompt label tied to the selected physical reference slot.

Audio is more dynamic. A voice set to **Automatic** can be bypassed when its bound Subject does not speak in the current Generation. The composer rebuilds the active logical `<Audio N>` order for that Generation so you do not have to renumber the prompt manually.

If a retained soundtrack from a reference Video is active, that embedded audio can also appear before standalone Audio references in the H3 audio-label order. The composer handles that mapping automatically.

> **IMPORTANT:** Do not manually rewrite generated `<Audio N>` labels just because the physical ComfyUI socket has a different zero-based name.

---

5. Subjects

A **Subject** is reusable visible content in the target video. It can be:

- a character or person;
- a creature or animal;
- a vehicle;
- an object or prop;
- an environment/location;
- another reusable visual element.

Give each Subject a clear name. In Ref2VA, the composer assigns it a stable `<Subject N>` label in the generated prompt.

Character / appearance source images

Attach one or more Pictures that define the character.

A normal character sheet can establish:

- identity;
- face and hair;
- body proportions;
- wardrobe;
- accessories;

- distinguishing details.

Try to keep source images internally consistent. More references are not automatically better.

Wardrobe-only reference

Use **Wardrobe only** when an image should transfer clothing without replacing the target identity, scene, lighting, or framing.

Custom contribution

Use a custom contribution when a reference has a narrow, specific job that is not covered by the standard presets.

---

6. Environments and Multiple Views

An Environment is treated as its own Subject.

You may select **one or more views of the same location**. When several images are attached, the composer frames them as complementary views of one stable environment rather than separate target frames.

For example, the output may say:

```
<Subject 3> is the apartment living room, the same location shown in
<Picture 3> and <Picture 4>. Treat these as complementary views of one
stable environment; preserve its architecture, layout, furnishings,
materials, colors, and lighting.
```

This is useful when one reference shows the couch side of a room and another shows the reverse angle.

Best practices for multi-view environments

- Use images that genuinely depict the **same place**.
- Keep architecture, furniture, materials, and lighting reasonably consistent between views.
- Use the Shot and Camera Builder to describe the target viewpoint. Do not expect the environment references themselves to act as target framing unless you intentionally create a frame anchor.
- If two source images strongly disagree about the layout, H3 may still have difficulty reconciling them.

---

7. Pictures: Subject Source vs. Frame Anchor

There is an important distinction between two kinds of image use.

Subject source image

If Picture 1 is simply the image that defines a character, attach it to that Subject. The composer references it inside the Subject definition.

Standalone Picture reference

Create a standalone Picture reference only when the image itself has a separate job, such as:

- opening frame;
- ending frame;
- composition anchor;
- storyboard/frame-planning anchor;
- wardrobe transfer;
- character scale reference;
- environment state reference.

This helps prevent every source image from being treated as a target composition.

---

8. Character Height and Scale

The composer supports a dedicated **Character Scale Reference** Picture.

This is useful when a simple composite shows two characters standing on the same floor line and you want H3 to preserve their relative height/body scale.

Choose:

- the connected Picture used for the scale diagram;
- the Subject shown on the left;
- the Subject shown on the right, if present.

The generated prompt limits that Picture to attributes such as:

- relative height;
- body scale;
- proportions;
- shared floor-plane contact.

It explicitly avoids using the diagram's left/right placement as target blocking or camera framing.

If you do not want to use another image, **Text-only height relationships** are available as an alternative.

---

9. Generations and Continuity

The composer can hold multiple Generations inside one project.

New Generation

A fresh Generation starts with a new Shot but carries forward useful scene continuity:

- base soundscape;
- music choice;
- carried appearance state for supported Subjects.

Generation-specific extra sounds are intentionally reset.

Duplicate Generation

Duplicate when you want the same Shots and settings as a starting point for a variation.

Appearance continuity

Supported character Subjects can have reusable **Appearance States**, for example:

- Base / character-sheet appearance;
- Jacket On;
- Jacket Off;
- Wet Clothes;
- Injured;
- Evening Wardrobe.

A later Generation can inherit the ending appearance from the previous one. You can also recall an earlier saved appearance state.

When editing a state that is already used historically, the composer is designed to protect earlier Generations by allowing the current Generation to fork a variation rather than silently rewriting the past.

---

10. Building a Shot

Each Shot card contains the information that actually happens in that continuous camera setup.

Environment / location

Choose the Environment Subject used in the Shot.

Subject placement and action before dialogue

Use the main action box for:

- opening placement;
- opening physical state;
- untimed action;
- an action that clearly occurs before dialogue.

Example:

```
<Subject 1> stands beside an open suitcase on the bed.
```

If all of the Shot's actions are timed with Timed Action Beats, this box can be left empty.

> **RULE OF THUMB:** Do not write the same action in both the normal action box and a Timed Action Beat. Give each instruction one home.

---

11. Timed Action Beats

Use **Timed Action Beats** when multiple actions need specific pacing inside one continuous Shot.

Example: two equal actions in a 10-second Shot

```
0.00 - 5.00
<Subject 1> folds the clothes and places them into the suitcase.

5.00 - 10.00
<Subject 1> closes the suitcase, stands, and carries it toward the door.
```

The compiler produces natural timeline language such as:

```
From 0.00 to 5.00 seconds into the shot, <Subject 1> folds the clothes...
From 5.00 to 10.00 seconds into the shot, <Subject 1> closes the suitcase...
```

Distribute evenly

Add the number of beats you want, then click **Distribute evenly**.

For a 10-second Shot:

- 2 beats -> 0-5 / 5-10
- 4 beats -> 0-2.5 / 2.5-5 / 5-7.5 / 7.5-10

You can adjust the values afterward.

Shot-relative timing

Timed Beat values are relative to the current Shot.

If Shot 2 begins six seconds into the Generation and lasts four seconds, you still enter its local beats as:

```
0 - 2 sec
2 - 4 sec
```

You do not need to mentally convert them to 6-8 and 8-10.

Timed Beats and dialogue

Timed Action Beats currently pace **actions**, while dialogue remains in Dialogue cards.

If a timed action happens while somebody speaks, write the action so simultaneous performance makes sense. If the dialogue needs its own exclusive time, leave room in the Shot rather than filling every second with incompatible actions.

Timed Beats and Camera Motion

If timed beats fill essentially the entire Shot and the camera is supposed to move at the same time, use a simultaneous Camera Motion timing such as:

Throughout the shot

Do not choose **After opening action** if you have already assigned timed action through the end of the Shot and expect the camera move to happen simultaneously. Prompt Check will warn about that chronology conflict.

---

12. Camera Builder

The Camera Builder is endpoint-first: describe where the Shot begins, where it should end, and then choose a physically compatible way to get there.

Step 1 - Start Frame

Typical controls include:

- framing / shot size;
- primary Subject or Environment;
- optional second Subject;
- horizontal viewpoint;
- camera height / vertical angle;
- composition;
- advanced framing notes.

Step 2 - End Frame

The End Frame can stay the same or independently change:

- shot size;
- framed Subject(s);
- viewpoint;
- camera height;
- composition.

This is especially useful for moves such as:

- Medium two-shot -> Medium close-up isolating Subject 1;
- Medium close-up Subject 1 -> wider two-shot revealing Subject 2;
- left three-quarter -> right three-quarter orbit.

Step 3 - Camera Motion

The composer analyzes the Start and End Frames and filters movements by physical compatibility.

Examples include:

- Push In / Pull Out;
- Truck Left / Right;
- Pan Left / Right;
- Tilt Up / Down;
- Pedestal Up / Down;
- Crane / Jib Move;

- Dolly Diagonally;
- Orbit / Arc;
- Tracking Shot;
- Zoom In / Out;
- Dolly Zoom;
- Roll;
- Static Shot;
- Custom Movement.

If an endpoint change makes the currently selected movement physically incompatible, the composer can clear or flag it rather than silently generating contradictory instructions.

Motion timing

Camera Motion can be placed:

- After opening action / before dialogue;
- During opening action;
- During dialogue;
- After dialogue;
- Throughout the shot.

Choose the timing that describes when the camera actually moves, not merely the one that sounds most general.

Advanced camera controls

Advanced view includes controls for items such as:

- movement range and speed;
- stabilization / handheld character;
- lens perspective;
- depth of field;
- focus behavior;
- rack focus;
- custom movement instructions.

The generated prompt uses natural camera language rather than exposing internal UI terms.

---

13. Dialogue and Voices

Add a **Dialogue** card when someone speaks.

For each line you can set:

- speaker Subject;
- delivery/performance direction;
- language;
- exact dialogue;

- voice reference behavior;
- off-screen / voice-over options;
- an action that happens after the line.

The composer generates H3 speaker IDs such as `(S1)` and dialogue tags such as:

```
<d>[English] I don't think I'm ready to open it.</d>
```

Bind a voice once

A voice Audio reference can be bound to a Subject. In normal use, leave the Dialogue card's voice selection on **Automatic from subject**.

If that Subject does not speak in the current Generation, the voice is automatically bypassed and does not consume a logical `<Audio N>` label.

This is particularly useful when a project contains several characters but only some of them speak in a given Generation.

Action after this line

Use **Action after this line** for reactions or behavior that must happen after the spoken dialogue.

Example:

```
She closes her lips, lowers the envelope to the counter, and waits in silence.
```

This keeps timeline order clear instead of putting a post-dialogue reaction into the opening-action box.

---

14. Video Reference Workflows

The composer includes guided setups for several common Ref2VA tasks.

Insert Subject Into Existing Footage

Use this when a generated Subject should be inserted into a source video while preserving the original plate/camera.

The guided workflow can describe:

- inserted Subject;
- source environment;
- placement/depth;
- occlusion;
- grounding/contact;
- source-audio handling.

Because the source video owns the camera in this workflow, Camera Builder settings are suppressed to avoid conflicting instructions.

Background Replacement + Relighting

Use this when a performer exists in source footage and you want to replace the background with a referenced Environment, then integrate the performer into the new lighting.

The workflow can preserve performance while specifying:

- replacement Environment;
- background removal type;
- spill cleanup;
- integration strength;
- relighting details;
- source-audio treatment.

Performance + Camera Transfer

Use this when a filmed performance provides acting, facial behavior, gesture, timing, and optionally camera movement to a generated target character.

This is performance transfer, not green-screen compositing.

---

15. Soundscape and Music

The composer separates scene sound from audience-only music.

Overall soundscape

Use this for sounds that belong to the world of the video, such as:

- room tone;
- traffic;
- wind;
- footsteps;
- clothing movement;
- doors;
- physical impacts;
- non-verbal human sounds.

Do not repeat the exact spoken dialogue here.

A base soundscape can carry into a new Generation, while **Additional sounds this Generation** can be used for clip-specific events.

Non-diegetic music

This is music audible to the audience but not the characters.

Describe concrete musical qualities such as:

- instrumentation;

- tempo;
- rhythm;
- dynamics.

Choose **No non-diegetic music** when there is no score.

Choose complete silence for the soundscape only when the entire video is genuinely silent.

---

16. Prompt Check

Prompt Check is the composer's built-in validator.

Error

An **Error** means the composer has found a blocking structural or mapping problem that should be fixed before copying the prompt.

Examples:

- invalid reference slot;
- timed beat outside the Shot duration;
- overlapping or zero-length timing that cannot be interpreted correctly;
- incompatible camera movement;
- missing required sound configuration.

Warning

A **Warning** means the prompt can still be generated, but the settings may conflict or produce unclear chronology.

Examples:

- Timed Action Beats fill the full Shot but the camera is told to move afterward;
- an End Frame changes but no Camera Motion is selected;
- a movement target looks stale after reframing a duplicated Shot.

Info

An **Info** message is advisory. It is not a failure.

Examples:

- H3 duration snapping to the 17n+5 frame grid;
- an unused automatic voice being bypassed;
- a `detailed_description` word-count reminder;
- an intentional gap between Timed Action Beats.

A clean Prompt Check means the composer does not see a structural conflict. It does **not** guarantee that the generative model will follow every instruction perfectly.

---

17. Duration and the H3 Frame Grid

The composer shows both your requested duration and H3's effective frame-grid duration.

For example, a requested 10-second target may snap to approximately:

```
243 frames = 10.13 sec at 24 fps
```

Timed Action Beats intentionally use the **creative target duration you entered**. If you ask for a 10-second Shot and distribute two beats evenly, you get 0-5 and 5-10 rather than awkward values based on the small frame-grid tail.

---

18. Built-in Examples

The composer includes example projects for learning and testing:

1. **Full reference scene -** two characters, two voices, environment, scale reference, two Shots, camera movement, soundscape.
2. **Insert subject -** add a generated subject to existing footage.
3. **Background replacement + relight -** replace a source background and integrate the performer.
4. **Performance + camera transfer -** transfer a filmed performance to a target character.
5. **Simple T2VA -** demonstrates a straightforward text-to-video Shot and Timed Action Beats.

A useful way to learn the composer is to open an example in a second browser window and use it as a reference while building your own project.

---

19. Basic vs. Advanced View

Basic view exposes the controls most users need for normal generations.

Advanced view reveals additional options for fine control, including more detailed reference, appearance, camera, focus, and workflow settings.

If you are unsure what an advanced option does, leave it at the default. The composer is designed so that unused controls do not need to be filled simply because they exist.

---

20. Saving and Sharing Projects

Automatic local save

The composer automatically saves its state in the browser's local storage for that version.

When a newer compatible version is opened, it can migrate supported older project state forward.

Save Project

Use the project-save control to export a JSON project file when you want a portable backup or need to move the setup to another computer/browser.

Restore Previous

Before destructive actions such as loading a guided preset or example over meaningful work, the composer keeps a previous-project backup that can be restored.

V5.11.1 also recognizes meaningful work contained only in Timed Action Beats or earlier Generations, so those projects are protected by the replacement warning.

---

21. Practical Prompt-Building Principles

The composer is built around a few simple principles:

Define something once where it belongs

A character's identity belongs in the Subject definition. A Shot describes what that character does now. Avoid re-explaining the entire character sheet in every Shot.

Give each reference one clear job

If an image is a character sheet, use it as a character source. If it is a scale diagram, use it for scale. If it is an exact frame, use it as a frame anchor.

Describe the timeline in playback order

Opening state -> action -> camera behavior -> dialogue -> reaction should read in the order the viewer experiences it, unless events intentionally happen at the same time.

Prefer concrete instructions over compiler explanations

The generated prompt should tell H3 what is visible, audible, moving, or changing - not explain how the application arrived at the instruction.

More text is not automatically better

The goal is high signal-to-noise. Add detail when it resolves ambiguity, preserves continuity, or controls an important result.

---

22. Troubleshooting

The wrong voice is used

Check:

1. The Audio reference is bound to the correct Subject.
2. The correct Subject is selected as the Dialogue speaker.

3. Voice override is either Automatic or intentionally set.
4. The connected Audio count matches ComfyUI.
5. Read Prompt Check to see whether another automatic voice was bypassed and the logical `<Audio N>` order changed.

Do not assume the number painted on a zero-based physical socket must be copied directly into the prompt.

A character appears when it should not

Make sure that Subject is not accidentally referenced in:

- the Shot action;
- Timed Action Beats;
- dialogue;
- Camera Builder target/foreground/second-subject controls;
- a frame anchor or workflow setting.

The composer tries to avoid globally mentioning unrelated characters because unnecessary references can encourage H3 to introduce them.

Two characters merge or swap identity

Use separate Subject definitions and clean reference images. When two defined Subjects are actually co-present, the composer can add concise identity-separation guidance.

Environment details change

First confirm that your environment reference views depict the same coherent location. If Prompt Check is clean but the generated room still shifts, remember that model, sampler, acceleration, quantization, and seed behavior can affect reference fidelity independently of prompt structure.

Timed Beat warning about camera movement

If the actions occupy the whole Shot and the camera is intended to move at the same time, choose **Throughout the shot** or another simultaneous camera timing.

Word-count information message

The official full-reference format generally expects a detailed generation description, but the composer does not add filler simply to reach a number. Treat the word-count message as an advisory check. Add concrete visual/timeline detail when the Shot genuinely needs it.

---

23. What the Composer Does Not Do

H3 Prompt Composer does not:

- run MiniMax H3;
- connect to the internet;
- upload reference media;

- change your ComfyUI workflow;
- choose your model checkpoint, sampler, step count, or acceleration settings;
- guarantee that H3 will follow every prompt instruction.

Its job is narrower: **build a clear, internally consistent H3 prompt from filmmaking-oriented controls and catch common structural mistakes before generation.**

---

24. Quick Reference

If you want to...	Use...
Define a reusable character	Subject
Define the room/location	Environment Subject
Use two angles of the same room	Multiple Environment source images
Lock relative character height	Character Scale Reference or Text-only height relationship
Make a second camera angle	New Shot
Perform a second action without a cut	Timed Action Beat
Split two actions evenly across 10 sec	Two Timed Beats + Distribute evenly
Move the camera while timed actions occur	Camera Motion -> Throughout the shot
Specify exact spoken words	Dialogue card
Preserve a character's voice	Voice Audio bound to that Subject
Put a reaction after speech	Action after this line
Change wardrobe later	Appearance State
Show the wardrobe change on camera	Appearance change during this Shot
Carry scene ambience forward	Base soundscape
Add one-off footsteps/impact	Additional sounds this Generation
Insert a character into source footage	Insert Subject workflow
Replace a background and relight	Background + Relight workflow
Transfer filmed acting	Performance + Camera Transfer workflow
Check the final structure	Prompt Check

Release

H3 Prompt Composer V5.11.1 is the audited sharing build of the Timed Action Beats release.

The application is self-contained and local. Keep the HTML file and this guide together when distributing it so new users have both the tool and the operating instructions.