

H3 Prompt Composer

User Guide

V5.12.4

Standalone local HTML tool for building structured MiniMax H3 prompts

Built for reference-heavy generation, continuity, video editing, performance transfer, and continuation workflows.

Contents

1. Quick Start
 2. The Three Timeline Levels
 3. Choosing a Mode
 4. Setting Up Ref2VA References
 5. Subjects and Source Media
 6. Environments and Multiple Views
 7. Standalone Picture Roles
 8. Height, Scale, and Placement
 9. Appearance and Wardrobe Continuity
 10. Generations and Continuity
 11. Building a Shot
 12. Timed Action Beats
 13. Camera Builder
 14. Dialogue and Audio
 15. Video Reference Workflows
 16. Soundscape and Music
 17. Prompt Check
 18. Duration and the H3 Frame Grid
 19. Built-in Examples
 20. Basic vs. Advanced View
 21. Saving and Sharing Projects
 22. Practical Prompt-Building Principles
 23. Troubleshooting
 24. What the Composer Does Not Do
 25. Quick Reference
- Release Notes

1. Quick Start

1. Open H3_Prompt_Composer_V5_12_4.html in a modern browser.
2. Choose the H3 mode you are using: T2VA, I2VA, FL2VA, L2VA, or Ref2VA.
3. Set the target duration and a short visual style description.
4. For Ref2VA, define only the Subjects and references that have a clear job. In Advanced view, match the connected Images, Videos, and Audio counts to your ComfyUI inputs.
5. Build the Generation from one or more Shots. Add another Shot only when the camera actually cuts or transitions to a new view.
6. Describe the opening state/action, optional Timed Action Beats, camera, dialogue, and audio.
7. Read Prompt Check. Fix Errors, review Warnings, and use Info messages as guidance.
8. Use Copy prompt and paste the result into the matching H3 prompt field in ComfyUI.

Best first move Open Load example and choose the workflow closest to what you want to do. Replace its references and actions one section at a time.

The Composer is not a media loader The HTML does not load the images, videos, or audio connected to ComfyUI. It describes how those already-connected inputs should be interpreted.

2. The Three Timeline Levels

Level	What it means	Create a new one when...
Generation	One complete H3 output/prompt. It has its own Shots, soundscape, music, active references, and appearance continuity.	You are making the next output clip.
Shot	One continuous camera view inside a Generation.	There is a real cut, transition, or new camera view.
Timed Action Beat	A timed action window inside one continuous Shot.	The camera does not cut, but Action A and Action B need controlled pacing.

Rule of thumb If the camera stays continuous, keep one Shot. Use Timed Action Beats to control what happens during different parts of that Shot.

3. Choosing a Mode

Mode	Use it when...	Main idea
T2VA	Generating from text without an image keyframe.	Build the audiovisual timeline from text.
I2VA	Picture 1 is the required opening	Begin at Picture 1 and develop

	frame.	forward.
FL2VA	Picture 1 is the first frame and Picture 2 is the last frame.	Describe a continuous path between the two frames.
L2VA	Picture 1 is the required final frame.	Build toward the final image.
Ref2VA	Using character, environment, Picture, Video, or Audio references; editing; transfer; continuation.	Define references once, then describe how the target video uses them.

Ref2VA outputs six sections: `subject_definitions`, `summary`, `retention_analysis`, `detailed_description`, `overall_soundscape`, and `non_diegetic_music`. The Composer formats these automatically.

4. Setting Up Ref2VA References

In Basic view, add Subjects and reference roles with the large workflow and add-reference controls. Advanced view exposes Connected ComfyUI inputs so you can set the physical image, video, and audio counts when needed.

Reference numbering

ComfyUI may show zero-based socket names such as `ref_image_0`. H3 prompt labels are one-based, such as `<Picture 1>`, `<Video 1>`, and `<Audio 1>`. Picture and Video labels stay tied to their physical slots. Audio labels are rebuilt per Generation from the active audio relationships.

Do not hand-renumber Audio labels A voice in Auto mode can be bypassed when its Subject does not speak. Retained Video soundtracks can also occupy an Audio label before standalone Audio inputs. The Identity and Voice Map shows the active mapping.

Give every reference one clear job

A clean Ref2VA setup is easier for both the user and the model to interpret. Avoid attaching the same image in several roles unless each role is intentional and necessary.

5. Subjects and Source Media

A Subject is reusable visible content: a character/person, creature/animal, vehicle, object/prop, Environment, or another reusable visual subject.

Subject / appearance source images

Select the Pictures that define the Subject. A normal character sheet can establish identity, face, hair, body proportions, wardrobe, accessories, and distinguishing features.

Optional subject / performance source video

This control appears only when Video inputs are configured. Use it when existing footage should help define the Subject's motion, performance, timing, wardrobe, or appearance. If you are using still images only, no Video source is needed.

Reference contribution

Choice	Use it for
Character sheet - identity and full appearance	Normal character/creature reference. Identity and appearance come from the attached Subject references.
Wardrobe only	Transfer clothing/accessories without replacing identity, pose, environment, lighting, or framing.
Environment identity and layout	Define a reusable location and spatial continuity.
Custom contribution	A narrow role that is not covered by the standard presets.

6. Environments and Multiple Views

An Environment is a reusable Subject representing a location. Attach every useful view of the same stable place to one Environment Subject. The Composer describes the images as complementary views rather than separate target frames.

- Use source views that genuinely depict the same location.
- Keep architecture, furniture, materials, colors, and lighting reasonably consistent.
- Use the Shot and Camera Builder to describe the target viewpoint. Environment references define the world and layout; they do not automatically own the final framing.
- An optional Environment source Video can be used when source footage of the same location should help define architecture, lighting, or continuity.

7. Standalone Picture Roles

A Picture that merely defines a Subject belongs inside that Subject. Add a standalone Picture only when the image itself has an extra job.

Picture role	What it controls
Composition anchor	Viewpoint, framing, and major subject placement for an assigned Shot.
Storyboard reference	Shot planning, blocking, action direction, framing intent, and planning timing. Annotations should not appear in the video.
Visual height / scale / placement	Height/body scale/floor contact globally, with optional approximate placement for one selected Shot.
Lighting integration only	Key/fill/rim relationship, reflected color, shadow density, exposure balance, local contrast, and atmospheric integration.
Wardrobe only	Clothing, materials, colors, accessories, fit, and layering for a selected Subject.
Environment state	The current condition of an existing Environment: weather, lights, moved props, damage, curtains, etc.

Exact first frame	Required visible composition at the start of the assigned Shot.
Exact keyframe	A required frame state reached naturally within the Shot.
Exact last frame	Required visible composition at the end of the Shot.
Exact edited keyframe	A specific edited target frame state at the intended point in the Shot.

Avoid accidental framing pressure Do not turn every useful reference image into a standalone Picture. Subject and Environment source images should stay inside their Subjects unless the image has a second, deliberate role.

8. Height, Scale, and Placement

Text-only height relationship

Use written height guidance when that is enough. It is the lightest and safest option because it adds no extra visual composition for H3 to reconcile.

Visual height / scale / placement reference

Use a deliberately simple scale diagram when you need stronger physical-size guidance. The guide can show one character or two characters on a shared floor line.

- **Height / scale only - safest:** Transfers height, body scale, proportions, and floor-plane information without using the guide as target blocking.
- **Height / scale + approximate placement:** Also transfers approximate blocking/depth for one selected Shot. This is useful for a rough Photoshop/After Effects composite showing how large and roughly where the inserted Subject should appear.

In both modes the prompt explicitly tells H3 not to reproduce the guide's pose, lighting, rough mask edges, image style, or exact composition.

9. Appearance and Wardrobe Continuity

Core idea Starting look for this Generation is the authoritative opening appearance. The fields directly below it always edit the look that the current Generation is actually using.

Starting look for this Generation

A fresh Generation inherits the carried-forward look from the Generation you create it from. If wardrobe changes off camera between clips, choose an existing saved look or click + New look for this Generation.

Creating and naming a new look

9. Click + New look for this Generation.
10. Type a meaningful Starting look name such as Evening Wardrobe or Jacket Off.
11. Press Enter or click Rename. The Starting look dropdown updates immediately.

12. Choose how the look is defined: Base character sheet, Text description, Dedicated appearance Picture, or a panel/figure in a multi-state sheet.
13. If using a Text description, describe only the wardrobe/temporary appearance that overrides the base character sheet. Identity, face, hair, and body proportions remain anchored by the Subject reference.

Saved looks library / historical editing

The Saved looks library manages reusable appearance states; it is not the selector for the current Generation. Editing a saved look does not make the current Generation use it. Select the desired look with Starting look for this Generation.

Historical safety is on by design. If a look is already used in an earlier Generation, editing it later normally forks a variation for the current Generation instead of silently rewriting history.

When the appearance changes on camera

Describe the visible event in the normal Action or Timed Action Beats. Use Appearance result / continuity after this Shot only to record the resulting look that should carry into later Shots and Generations.

Action: <Subject 1> removes her jacket and drapes it over the chair.
Appearance after this Shot: Jacket Off

If the action already makes the result obvious, the Composer does not need to add a redundant sentence such as "she has no jacket at the end." A later Shot can receive a brief continuity cue that the Subject continues in the Jacket Off appearance.

10. Generations and Continuity

Control	What it does
+ New generation	Creates a fresh output with a new Shot while carrying appropriate continuity such as ending appearance, base soundscape, and music. Workflow-specific edit references do not leak into unrelated fresh Generations.
Duplicate	Copies the current Generation, including its Shots and workflow-specific settings, so you can create a variation.
Appearance carry-forward	The final tracked appearance state becomes the starting look for a newly created Generation.
Base soundscape	Carries forward to fresh Generations. Generation-specific extra sounds reset.
Music	Carries forward unless you change it.

Fresh vs Duplicate Use New generation for the next clip in a sequence. Use Duplicate when you want another version of the same setup.

11. Building a Shot

Each Shot represents one continuous camera view. The main Shot controls establish the Environment, opening subject placement/action, camera, dialogue, audio, and optional continuity changes.

Opening subject placement / action

Use this field for the opening placement/state and any untimed action. If dialogue is present, a reaction that must happen after a line belongs in that Dialogue card's Action after this line field.

<Subject 1> stands beside the open suitcase, facing the doorway.

If all physical action is already described by Timed Action Beats, the opening action field can be brief or empty. Do not write the same action twice.

12. Timed Action Beats

Timed Action Beats control pacing inside one continuous Shot. Their times are relative to the Shot, not the whole Generation.

0.00 - 5.00 sec <Subject 1> folds the clothes and places them into the suitcase.
5.00 - 10.00 sec <Subject 1> closes the suitcase, stands, and carries it toward the door.

Distribute evenly

Use Distribute evenly when you want equal time for each beat. A 10-second Shot with two beats becomes 0-5 / 5-10; four beats become 0-2.5 / 2.5-5 / 5-7.5 / 7.5-10.

Timed Beats with dialogue and camera

Actions can overlap dialogue when that is physically intended. If timed beats occupy the whole Shot and the camera also moves, choose a simultaneous camera timing such as Throughout the shot rather than scheduling the move after the actions end.

13. Camera Builder

The Camera Builder is endpoint-first: define where the Shot begins, where it should end, and then choose a physically compatible movement between those endpoints.

1 - Start Frame

- Framing / shot size
- Primary Subject or Environment
- Optional second Subject
- Horizontal viewpoint
- Camera height / vertical angle
- Composition
- Optional advanced framing notes

2 - End Frame

Leave controls as Same as Start unless the shot must finish differently. You can independently change shot size, framed Subject(s), viewpoint, height, composition, foreground/POV relationships, and other supported endpoint attributes.

3 - Camera Motion

The Composer analyzes the endpoints and filters motion choices by physical compatibility. Available moves include push/pull, truck, pan, tilt, pedestal, crane/jib, diagonal dolly, orbit/arc, tracking, zoom, dolly zoom, roll, static, and custom movement.

Motion timing and advanced controls

- After opening action / before dialogue
- During opening action
- During dialogue
- After dialogue
- Throughout the shot
- Advanced: movement range/speed, stabilization, lens perspective, depth of field, focus behavior, rack focus, and custom camera instruction override.

Source-video edit workflows When a direct edit preset preserves the source camera frame-for-frame, Camera Builder is intentionally ignored so it cannot fight the source footage.

14. Dialogue and Audio

Dialogue cards

Add a Dialogue card only when someone speaks. Choose the speaker, delivery direction, language, exact spoken words, optional voice behavior, off-screen/voice-over state, and an action that happens after the line.

Voice binding

Bind a Voice Audio reference to a Subject once. In normal use, leave the Dialogue card on Automatic from subject. If that Subject does not speak in the current Generation, the voice is bypassed and does not consume an H3 <Audio N> label.

Audio purposes

Purpose	Behavior
Voice timbre	Guides the generated speaker voice. Written dialogue controls what is said.
Exact performed dialogue	Reuses the recorded words, voice, timing, cadence, emotion, pauses, and pronunciation.
Exact scene soundtrack	Copies the source soundtrack relationship.
Reuse source audio	Preserves the source audio, optionally allowing requested

	additions.
Music style	Uses the clip as a music-style reference.
Sound texture	Uses the clip as a sound/SFX texture reference.

Non-voice audio uses Active / Bypass behavior; voice-oriented references can use Auto / Always / Off as appropriate.

15. Video Reference Workflows

The four guided workflows are the easiest way to build video-reference prompts. They automatically apply the source-ownership and retention language needed for each job.

A. Insert a Subject into existing footage

Use untouched source footage as Video 1 and the inserted Subject's appearance image as Picture 1. Video 1 owns the existing plate, timing, camera, parallax, environment, lighting, and motion blur. There is no redundant Environment Subject for the source plate.

- Choose foreground, midground, or background depth.
- Describe the inserted Subject's main path/blocking in What the inserted subject does.
- Optionally preserve source audio unchanged, preserve it and add edit sounds, or replace/generate new sound.
- Use existing objects as occluders, ground feet/wheels, add contact shadow, and match lighting/focus/grain/motion blur.
- For stronger scale or blocking control, add a Visual height / scale / placement Picture using a rough composite.

B. Background replacement + relighting

Use performer footage as Video 1 and the replacement Environment as a Picture-backed Environment Subject. Choose the original background type (green, blue, neutral, dark, rough temporary composite, or existing location) and the desired relighting strength.

- Replace background + relight in one pass when the source still needs background replacement.
- Relight existing composite only when the background is already correct and only the performer illumination/shading needs integration. This is useful for a composite created in After Effects or for using the output of a first H3 pass as the next source.
- Optionally add a Lighting / Integration Reference Picture showing a preferred relit/composited still. It transfers lighting relationships only, not pose, framing, placement, or background geometry.

C. Transfer performance or camera

Use a target character sheet as Picture 1 and a short filmed performance/camera reference as Video 1. Select exactly which attributes should transfer: facial expression, mouth performance, head movement, eyeline, upper-body gesture, full-body motion, and/or camera movement.

Dialogue/audio is optional. Choose No dialogue / no audio reference for movement/acting/camera transfer with no spoken line. Exact performed dialogue can reuse the recorded speech signal when needed.

D. Continue an existing video

Use Video 1 as a temporal handoff. The source clip ends; the target begins immediately afterward. This is not a frame-for-frame edit.

- Continue Subject state, movement, and screen direction.
- Continue camera direction, speed, height, and framing momentum.
- Preserve Environment, lighting, atmosphere, and spatial continuity.
- Avoid replaying, restaging, or resetting the ending action.
- Use Handoff notes for details such as foot position, body momentum, camera drift, eyeline, or a prop state at the exact handoff.

Begin immediately after the final moment of <Video 1>, with no cut, pose reset, camera reset, or repeated action.

16. Soundscape and Music

Soundscape status

Choice	Meaning
Not configured	Prompt Check reports an Error until the scene sound is defined or silence is explicitly chosen.
Describe soundscape	Use the inherited base soundscape plus optional sounds specific to this Generation.
No scene sound / silence	No dialogue, ambience, or physical scene sound. Non-diegetic music is controlled separately.

Base soundscape

Use this for persistent room tone or ambience that should carry between Generations: distant traffic, HVAC hum, wind, birds, machinery, etc.

Additional sounds this Generation

Use this for clip-specific physical sounds such as footsteps, a door slam, cooking, impacts, cloth movement, or inserted-subject effects. These do not automatically carry into a fresh Generation.

Non-diegetic music

Describe audience-only score in concrete musical terms: instrumentation, tempo, rhythm, density, and dynamics. Choose No non-diegetic music when there is no score.

17. Prompt Check

Status	Meaning	What to do
Error	A blocking structural, mapping, timing, or required-configuration problem.	Fix before generating.

Warning	The prompt can build, but a chronology/reference/camera relationship deserves review.	Change it if the warning describes something you did not intend.
Info	Helpful context such as frame-grid snapping, word-count guidance, automatic voice bypass, or workflow advice.	Usually no action is required.

A clean Prompt Check means the Composer does not see a structural conflict. It does not guarantee that the generative model will follow every instruction perfectly.

18. Duration and the H3 Frame Grid

The Composer displays the requested creative duration and the effective 24 fps frame-grid duration used by its H3 alignment logic. Timed Action Beats continue to use the clean creative duration you entered, so a requested 10-second Shot can still divide naturally into 0-5 and 5-10.

19. Built-in Examples

Example	What it demonstrates
Full reference scene	Two characters, two voices, multi-view Environment, visual scale reference, two Shots, dialogue, camera moves, soundscape.
Insert a subject into existing footage	Source plate ownership, placement, occlusion, grounding, integration, source audio.
Background replacement + relight	Replacement Environment, background-removal type, directional relighting, preserved performance.
Transfer performance + camera	Filmed acting/camera transfer with exact performed dialogue.
Continue an existing video	Temporal handoff from the final source moment into new action.
Simple T2VA shot	Text-only action, Timed Action Beats, Camera Builder, soundscape, music.

20. Basic vs. Advanced View

Basic view exposes the controls most people need. Advanced view reveals lower-level reference/task controls, Connected ComfyUI input counts, retention overrides, more appearance options, and advanced camera/focus controls.

Default advice Start in Basic. Switch to Advanced only when you have a specific reason to change a lower-level setting.

21. Saving and Sharing Projects

- **Auto-saved locally:** Current work is stored in browser local storage for that browser/profile.
- **Save .json:** Exports an editable project file that can be moved to another browser/computer.
- **Open:** Loads a saved Composer JSON project.
- **Restore previous:** Restores the safety backup made before destructive resets/preset/example changes.
- **Copy prompt:** Copies the active Generation prompt.
- **Copy all generations:** Copies every Generation in sequence for planning a multi-clip project.

Browser local storage does not travel with the HTML file. Export the JSON when you need a portable editable setup.

22. Practical Prompt-Building Principles

- Define reusable identity once where it belongs. Do not re-explain the entire character sheet in every Shot.
- Give every reference one clear job.
- Describe the timeline in playback order: opening state -> action -> camera behavior -> dialogue -> reaction, unless events intentionally overlap.
- Use Shots for cuts; use Timed Action Beats for pacing inside one continuous camera setup.
- Let Camera Builder own camera language whenever possible.
- Prefer concrete visible/audible instructions over explanations of how the Composer works.
- More text is not automatically better. Add detail when it resolves ambiguity, preserves continuity, or controls an important result.

23. Troubleshooting

Symptom	First things to check
Wrong voice	Confirm the Audio reference is bound to the correct Subject; confirm that Subject speaks in this Generation; inspect the Identity and Voice Map and Prompt Check.
Character appears unexpectedly	Check Shot action, Timed Beats, dialogue, Camera Builder targets, frame anchors, and workflow settings for accidental references.
Two characters merge/swap	Use distinct Subject references and keep the distinct-subject guard enabled when the characters are truly co-present.
Environment changes	Confirm all Environment views depict the same coherent place. Reduce conflicting references.
Wardrobe from an older look keeps returning	Check Starting look for this Generation. A text field in the Saved looks library does not select that look for the current Generation.
On-camera wardrobe change does not carry forward	Track Appearance after this Shot and make sure the visible change itself is described in the Action or Timed Action Beats.
Second action is rushed	Split the action into Timed Action Beats.

Camera timing warning	If actions fill the Shot and camera motion is simultaneous, choose Throughout the shot or another simultaneous timing.
Inserted Subject scale/position drifts	Try a Visual height / scale / placement guide and assign approximate placement to the affected Shot.
Relight looks like a color tint	Use directional lighting notes and optionally a Lighting / Integration Reference still.
Continuation repeats the ending	Enable the no-replay/reset behavior and add a precise Handoff note describing the final source state.
Prompt feels too long	Remove repeated statements first. Define each reference once and describe the actual timeline.

24. What the Composer Does Not Do

- Run MiniMax H3.
- Connect to the internet.
- Upload your reference media.
- Change your ComfyUI workflow.
- Choose your checkpoint, sampler, step count, quantization, or acceleration settings.
- Guarantee that H3 will follow every instruction perfectly.

Its job is narrower: build a clear, internally consistent H3 prompt from filmmaking-oriented controls and catch common structural mistakes before generation.

25. Quick Reference

If you want to...	Use...
Define a reusable character	Subject
Define a reusable location	Environment Subject
Use multiple views of one room	Multiple Environment source images
Lock relative character height	Text-only height relationship or Visual height/scale reference
Guide approximate scale + placement in one Shot	Visual height / scale / placement reference
Transfer clothing only	Wardrobe Picture role
Show a preferred relighting result	Lighting integration Picture role
Make a second camera angle	New Shot
Pace multiple actions without a cut	Timed Action Beats
Specify exact spoken words	Dialogue card
Preserve voice timbre	Voice Audio bound to the Subject

Reuse exact performed speech	Exact performed dialogue Audio
Change wardrobe between Generations	Starting look for this Generation
Show a wardrobe change on camera	Action/Timed Beat + Appearance after this Shot
Insert a character into source footage	Insert Subject workflow
Replace a background and relight	Background + Relight workflow
Relight an already-composited video	Relight existing composite only
Transfer filmed acting/camera	Performance + Camera Transfer
Transfer movement with no dialogue	Performance Transfer + No dialogue / no audio reference
Continue from the end of a source clip	Continue an existing video
Check final structure	Prompt Check

Release Notes

H3 Prompt Composer V5.12.4 is the current audited sharing build. The V5.12 series added cleaner source-video editing, scale + placement guidance, lighting-integration references, dialogue-free performance transfer, a guided continuation workflow, and a redesigned appearance/wardrobe continuity system. V5.12.4 also audited UI labels/helper text so visible instructions match the controls and reference roles currently on screen.

The application remains a single standalone local HTML file. Keep this guide with the matching V5.12.4 HTML when distributing the tool.