

H3 PROMPT COMPOSER

V5.14.0 User Guide

CURRENT SHARING BUILD

A practical reference for building structured MiniMax H3 prompts from filmmaking-oriented controls: Subjects, environments, multi-reference Pictures, Shots, camera behavior, dialogue, continuity, source-video workflows, and audio.

V5.14.0 audit-hardening release: The multi-Subject scale/placement workflow is preserved, while built-in example enums, generated grammar, migration self-healing, Prompt Check coverage, detailed-description word counting, and documentation have been tightened after an independent audit.

Contents

1. Quick Start
2. The Three Timeline Levels
3. Choosing a Mode
4. Connected ComfyUI Inputs
5. Subjects
6. Environments and Multiple Views
7. Standalone Picture Roles
8. Visual Height / Scale / Placement
9. Appearance and Wardrobe Continuity
10. Building a Shot
11. Timed Action Beats
12. Camera Builder
13. Dialogue and Audio
14. Video Reference Workflows
15. Video Continuation
16. Soundscape and Music
17. Prompt Check
18. Generations and Project Organization
19. Saving / Migration
20. Practical Prompt-Building Principles
21. Troubleshooting
22. Privacy and Security
23. Quick Reference

1. Quick Start

1. Open the HTML file in a modern browser.
2. Choose T2VA, I2VA, FL2VA, L2VA, or Ref2VA.
3. Set the target duration and visual style.
4. For Ref2VA, set **Connected ComfyUI inputs** to match the physical Image / Video / Audio reference sockets you connected.
5. Add reusable Subjects and Environments. Attach only the Pictures/Videos that define each one.
6. Add standalone Pictures only when they have a separate job: framing, storyboard, wardrobe, environment state, scale/placement, lighting integration, or an exact frame anchor.
7. Build each Generation from one or more Shots. Add Timed Action Beats, Camera Builder controls, dialogue, and sound as needed.
8. Read Prompt Check. Fix red errors, review amber warnings, and treat blue information messages as guidance.
9. Copy the generated prompt into the matching H3 conditioning node in ComfyUI.

Start simple. The Composer is designed around one-reference-one-job prompting. Add a control only when it solves a real ambiguity or continuity problem.

2. Timeline Levels and Modes

Generation

One complete H3 output/prompt. A fresh Generation starts the next clip while carrying useful continuity such as current appearance, base soundscape, and music.

Shot

One continuous camera setup. Create a new Shot only when the camera cuts or transitions to another view.

Timed Action Beat


A timed action window inside one continuous Shot. Use beats to pace multiple actions without inventing fake cuts.

Mode guide

Mode	Use it when...	Main relationship
T2VA	Text only	Build the audiovisual timeline from text.
I2VA	Exact first image	Picture 1 is the first frame.
FL2VA	Known first and last images	Travel continuously from Picture 1 to Picture 2.
L2VA	Known final image	Build toward Picture 1 as the ending frame.
Ref2VA	Characters, environments, voices, video references, edits	Define references once, then describe how each contributes.

3. Connected ComfyUI Inputs

In Ref2VA, **Images / Videos / Audio** are the physical reference slots you actually connected in ComfyUI. V5.14.0 keeps these controls visible in Basic view because they are fundamental setup, not an expert-only option.



If you set Images to 9, the Composer can offer Picture 1 through Picture 9 in Subject, Environment, and standalone Picture controls. The count does **not** mean you need nine Subject cards. Current H3 Ref2VA limits are up to 9 images, 3 videos, 3 audio clips, and 12 mixed reference files; individual audio/video references are 2-15 seconds with 15 seconds total per media type. H3 output duration is currently 4-15 seconds.

Important: Subject numbers are separate. One Environment Subject can use multiple Picture inputs when those images are complementary views of the same location.

4. Subjects and Environments

Subjects

A Subject is a reusable visible thing: character/person, creature/animal, vehicle, object/prop, Environment/location, or another visual element. Give it a clear name and select the Picture(s) that define identity/appearance.

For human characters, it is often useful to keep permanent physical identity in **Defining visual details** and use the Appearance system for wardrobe that may change.

Environments

An Environment is also a Subject, but its job is location identity and layout. If two or more Pictures show the same place from different angles, select all of them inside the same Environment card. H3 is told that they are complementary views of one stable location.

Environment references define the world, not the shot. The Shot and Camera Builder remain authoritative for the generated viewpoint and framing.

5. Standalone Picture Roles

A Picture attached inside a Subject defines that Subject. Create a standalone Picture entry only when the image itself has a separate job.

Role	What it contributes
Composition anchor	Viewpoint, framing, major placement for an assigned Shot.
Storyboard reference	Planning guidance: shot order, blocking, action direction, timing.
Wardrobe reference	Clothing/accessories only for a chosen Subject.
Environment state	Current location condition: weather, damage, lights, moved props, etc.
Visual scale / placement	Relative physical size and optionally approximate spatial placement.
Lighting integration	Lighting relationship and foreground/background integration only.
First / key / last / edited frame	Concrete frame anchor at a defined point in the Shot.

Keeping these roles separate prevents every image from accidentally becoming a target composition.

6. Visual Height / Scale / Placement

The current Composer uses an ordered multi-Subject list instead of the older two-Subject “left character / right character” control.

Four-subject scale / placement guide

Picture

<Picture 5>

auto definition

What does this image control?

Visual height / scale / placement – attribute guide

This Picture is an attribute guide, not an exact frame. Scale-only mode transfers relative physical size, body/object scale, proportions, and floor-plane information globally. The ordered Subject list identifies the figures without imposing target blocking. Scale + placement additionally transfers approximate left-to-right order, spacing, depth, and blocking for one chosen Shot. The prompt explicitly tells H3 not to reproduce the guide's pose, lighting, rough composite, or exact composition.

Short name

Four-subject scale / placement guide

Use a deliberately utilitarian scale diagram for scale-only work. A rough composite can also be used when you intentionally want approximate placement in one Shot. In either mode this Picture is guidance, never an exact target frame.

Connected Picture used for this scale / placement reference

Picture 5

Choose the actual ComfyUI Picture input containing the scale diagram or rough placement composite.

Subjects shown in this reference – from left to right

1<Subject 1> – Maya

2<Subject 2> – Eli

3<Subject 4> – T-Rex

4<Subject 5> – Jeep

+ Add Subject

List the visible Subjects in the same left-to-right order they appear in the guide. The order identifies which figure is which. In scale-only mode it does not control target blocking; in scale + placement mode it also guides approximate left-to-right placement for the selected Shot.

Use this image for this Generation

Height / scale + approximate placement

Placement applies to

Shot 1

Additional scale / placement notes (optional)

Optional: emphasize a specific depth, spacing, or floor-contact relationship; ignore pose, lighting.

Subjects shown in this reference - from left to right

Add the visible Subjects and arrange them in the same order they appear in the image. The list can contain people, creatures, vehicles, or props. Environments are excluded because this control describes object/character scale and placement within an Environment.

Within this guide, from left to right, the subjects shown are <Subject 1> Jane, <Subject 2> John, <Subject 4> T-Rex, and <Subject 3> Jeep.

Subject labels are stable IDs, not positions. A valid left-to-right mapping may therefore list Subject 4 before Subject 3; the ordered list tells H3 which stable label matches each visible figure.

H3 Prompt Composer V5.14.0

7. Scale-only vs. Scale + Placement

Height / scale only - safest

The ordered list identifies which figure is which, but the left-to-right arrangement does **not** define target blocking. The Composer transfers only relative physical size, body/object scale, proportions, and shared floor-plane contact.

Height / scale + approximate placement

The guide also contributes approximate left-to-right order, spacing, depth, floor contact, and screen placement for one selected Shot.

Still not an exact frame. Even in placement mode, the prompt explicitly rejects the guide's pose, wardrobe, lighting, image style, and rough-composite artifacts. Camera Builder and the Shot remain responsible for the actual shot design.

Migration

Older projects that stored a left Subject and optional right Subject automatically migrate those two Subjects into the equivalent ordered list when opened in V5.14.0. V5.14.0 also sanitizes invalid saved camera-enum values during migration.

8. Appearance and Wardrobe Continuity

Starting look for this Generation

This is the authoritative opening appearance. If the character uses the Base character sheet look, the source sheet remains authoritative. For a new wardrobe, use **+ New look for this Generation** and define the temporary wardrobe/appearance.

Rename

Type a new look name and press Enter or click Rename. The selector updates immediately. Historical safety prevents a later Generation from silently renaming/redefining earlier continuity unless you intentionally edit a shared state everywhere.

Visible wardrobe changes

Describe the action where it belongs: normal Action or Timed Action Beats. Then use **Appearance after this Shot** to record the resulting state for later Shots/Generations. The Composer does not redundantly narrate an obvious result such as “the jacket is now off” when the action already says she removes it.

▼ Appearance for this Generation

This section controls what H3 should see at the START of the current Generation. The fields directly below always belong to the selected Starting look. A fresh Generation inherits the carried-forward look from the Generation you create it from; you can change it here when wardrobe changes off camera between clips.

Starting look for this Generation

Evening Wardrobe ▼

This Generation is explicitly using the selected saved look. To reuse any older look, simply choose it here; a separate “Recall” control is not needed.

+ New look for this Generation

Starts as: Evening Wardrobe

Carries forward unchanged as Evening Wardrobe.

Starting look name

Evening Wardrobe

Type a new name; then press Enter or click Rename. The Starting look dropdown above updates immediately when the rename is committed.

Rename

How is this Starting look defined?

Text description — override wardrobe / temporary ▼

Starting wardrobe / temporary appearance used in this Generation

Black T-shirt and blue jeans, no jacket. Hair, face, and body proportions remain unchanged.

This text belongs to the selected Starting look above. It overrides conflicting wardrobe or temporary-appearance details from the base character sheet while identity, face, hair, and body proportions remain anchored by the Subject reference.

► Saved looks library / historical editing

9. Building a Shot and Timed Action Beats

Environment / location

Select which Environment Subject this Shot takes place in. A multi-Generation project may contain several Environments; each Shot chooses the one it needs.

Opening subject placement / action

Use the main Action field for opening state, placement, and untimed action. If several actions need specific pacing inside one continuous Shot, use Timed Action Beats.

Timed Action Beats

Beat times are Shot-relative. A 10-second continuous Shot can use 0-5 sec for one action and 5-10 sec for another without creating a cut.

One instruction, one home. Do not repeat the same physical action in the normal Action box and in a Timed Action Beat.

10. Camera Builder

The Camera Builder is endpoint-first: define where the camera starts, where it ends, then choose a physically compatible way to travel between them.

Camera Builder

Start: Medium shot · Maya · Front three-quarter — viewing more of subject's left side — End: Me...

1

START FRAME

Framing / shot size

Medium shot

Subject or Environment being framed

<Subject 1> — Maya

Second subject in frame — optional

— none / single-subject framing

Choose a second character here to make the selected shot size a two-person framing, such as a medium two-shot.

Horizontal viewpoint

Front three-quarter — viewing m

Camera height and vertical angle

Eye level, looking level

Composition — optional

Left third

2

END FRAME

Every blank End control inherits Start if a distinct endpoint is defined. If all End controls remain on "Same as Start," Step 3 can still use a loose movement range without forcing an exact ending frame.

Ending framing / shot size

Medium close-up

Subject or Environment being framed

Same as Start — <Subject 1> — I

Second subject in frame — optional

Same as Start — none / single-si

Keep the same second subject, switch to another, or choose None to resolve from two people to a single subject.

Horizontal viewpoint

Same as Start — Front three-qua

Camera height and vertical angle

Same as Start — Eye level, lookir

Composition — optional

Same as Start — Left third

3

CAMERA MOTION

Endpoint analysis: tighter ending framing. Choose how the camera physically gets from Frame 1 to Frame 2.

Movement

Push In

Recommended: Push In naturally connects the wider Start Frame to the tighter End Frame.

Movement target

<Subject 1> — Maya

Leave this on Automatic for normal subject-focused moves. Choose a different reference only when the camera is intentionally anchored to something other than the Start Frame target.

Speed

Slow

Motion timing

During dialogue

Controls when the movement happens relative to the shot action and dialogue; it changes prompt order instead of adding redundant timing instructions.

Additional movement instruction — optional

reveal <Subject 2> in the doorw.

Describe only something not already established by the End Frame.

► Stability, lens, and focus

CAMERA INSTRUCTION PREVIEW

A medium shot frames <Subject 1> from the waist up, with the camera at a front three-quarter angle favoring <Subject 1>'s left side, positioned on the left third of the frame. The camera is at eye level. During the dialogue, the camera slowly pushes in toward <Subject 1>, tightening to a medium close-up, with smooth stabilized movement.

Start Frame	End Frame	Camera Motion
Framing, target, optional second subject, viewpoint, camera height, composition.	Keep Same as Start or define a new framing/viewpoint/composition.	Static, Push/Pull, Pan, Truck, Tilt, Crane/Jib, Orbit/Arc, Tracking, Zoom, Dolly Zoom, etc.
Advanced optics controls include stabilization/handheld character, lens perspective, depth of field, focus behavior, rack focus, and custom movement text.		

11. Dialogue, Voices, Soundscape, and Music

Dialogue

Each Dialogue card defines speaker, delivery, language, exact spoken words, optional voice behavior, and optional action after the line.

(S1) Maya

Speaker

<Subject 1> – Maya

Voice: <Audio 1> – applied automatically from the subject binding.

quietly, with a hesitant pause before the final words

English

Spoken words only

I don't think I'm ready to open it.

► Action after this line – optional

► Advanced dialogue settings

Voice references

Bind a voice Audio reference to one Subject. Auto mode sends it only when that Subject speaks in the active Generation, allowing logical Audio labels to be rebuilt without wasting unused voice references.

Soundscape

Base soundscape carries forward into a fresh Generation. Additional sounds this Generation are clip-specific and reset.

Music

Non-diegetic music is audience-only score and is controlled separately from scene sound/silence.

12. Guided Video Reference Workflows

Insert Subject Into Existing Footage

Video 1 is the authoritative source plate. Preserve source timing, camera, parallax, environment, lighting, and unaffected objects; add only the new Subject and required interactions. Multi-Subject scale/placement guides can be used when a rough composite communicates size and blocking.

Background Replacement + Relighting

Preserve performer timing/framing, replace the source background, remove spill/edge contamination as appropriate, and rebuild directional lighting from the replacement Environment. A Lighting / Integration Reference can show a preferred relit still without transferring its pose or composition.

Performance + Camera Transfer

Transfer selected facial acting, mouth motion, head/eyeline, upper/full-body gesture, and/or camera movement while target Subject references remain authoritative for identity. Dialogue/audio is optional.

Choose the easiest starting point

Insert a subject into existing footage
Preserve the original source plate and camera frame by frame, then add a character, creature, vehicle, or prop with matching scale, occlusion, shadows, focus, and motion blur.

Replace a background and relight the performer
Use a green screen, blue screen, plain wall, neutral studio background, or existing composite. Replace the background with an image and transfer that environment's lighting onto the performer.

Transfer a filmed performance or camera move
Generate a referenced character while following facial acting, mouth performance, gesture, body motion, optional dialogue/audio, or camera movement from a short reference video.

Continue an existing video
Start immediately after the end of Video 1 while carrying forward subject motion, screen direction, camera momentum, environment, lighting, and spatial continuity.

Or build a general reference prompt

- 1 Add each reusable character, creature, vehicle, or prop as a **Subject**, and attach its appearance image or source video inside that Subject.
- 2 Add a location as an **Environment** and attach every useful view. These references define the world and layout, not the final camera framing.
- 3 Add a **Voice** and choose voice timbre or exact performed dialogue. In Auto mode, only voices used by the active Generation are sent to H3.
- 4 Add a standalone **Picture** only when the image

13. Video Continuation

Continuation is different from editing. A source-video edit modifies existing Video 1 frames. Continuation treats Video 1's final moment as a temporal/spatial handoff and creates new action after it.

Reference workflow: Continue an existing video

Leave preset

Connect the clip you want to continue as **Video 1**. The final moment of that video becomes the handoff into the new generation. This is not a frame-for-frame edit: the source ends, then the target video continues forward from its final state. Add Character / Subject references below only when you want stronger identity control.

<Video 1> continuation handoff

Preserve the ending state, then create new motion after it rather than replaying the source.

Continue subject state, movement, and screen direction

☒

Continue camera direction, speed, height, and framing momentum

☒

Preserve environment, lighting, atmosphere, and spatial continuity

☒

Do not replay, restage, or reset the ending action

☒

Handoff notes — optional

The clip ends with her left foot forward and the camera tracking backward; keep that stride and camera momentum continuous before she slows

Describe only the new action that happens after the source ends in the Shot action box. Camera Builder can be used when the continued camera deliberately evolves after the handoff.

You can preserve subject state/momentum/screen direction, camera direction/speed/height/framing momentum, and Environment/lighting/spatial continuity. The prompt explicitly says not to replay or restage the source ending.

Useful handoff notes: mention a specific ending foot position, gesture, gaze, vehicle trajectory, or camera direction only when it materially helps the transition.

14. Prompt Check

Prompt Check is the Composer’s preflight validator.

Duration snaps up to 243 frames = 10.13s on the 17n+5 frame grid.

detailed_description is about 253 words; generation tasks normally run 350–500 when the scene complexity supports it.

No blocking issues found. The prompt structure and mappings look ready.

Status	Meaning
Error	Blocking structural/mapping problem; fix before copying.
Warning	The prompt can build, but timing/reference/camera logic deserves review.
Info	Advisory guidance such as duration snapping, bypassed voice, or word-count context.

For multi-Subject scale guides, Prompt Check verifies that at least one Subject is listed, every entry maps to a valid non-Environment Subject, no Subject is duplicated, and placement mode points to a valid Shot. V5.14.0 also flags speaker IDs such as (S1) if they are manually typed into a custom retention note.

15. Generations and Project Organization

+ **New generation** creates a fresh output while keeping the project's reusable references available and carrying useful scene continuity such as appearance, base soundscape, and music. Each Shot can select a different Environment from the existing Environment Subject library.

Duplicate copies the current Generation when you want a variation of the same setup.

Project boundary suggestion: one coherent sequence/reference package usually works well as one Composer project, even if it moves through several related Environments. Start a new project when the reference package changes substantially, the project becomes unwieldy, or you need a clean organizational boundary.

16. Saving, Migration, and Practical Principles

Saving

- **Auto-saved locally** stores the current state in browser localStorage.
- **Save .json** creates a portable editable project file.
- **Open** loads a saved project JSON.
- **Restore previous** recovers the safety backup made before destructive resets/presets.

Migration

V5.14.0 recognizes supported older local-storage/project formats. V5.12.x two-Subject scale references migrate into the ordered multi-Subject list. Migration and Auto-tidy now sanitize unknown camera enum values, and Auto-tidy deduplicates scale-subject mappings while preserving first-seen order.

Principles

- Define reusable identity once.
- Give every reference one clear job.
- Describe the timeline in playback order.
- Let Camera Builder own camera language when possible.
- Prefer concrete visible/audible instructions over compiler explanations.
- More text is not automatically better.

17. Troubleshooting and Quick Reference

Symptom	First things to check
Another Picture is not selectable	Increase Connected ComfyUI Images count to match the physical reference inputs.
Two environment angles look like separate places	Put both Pictures inside one Environment Subject and confirm they genuinely depict the same location.
Scale guide maps the wrong figure	Check the ordered “Subjects shown in this reference - from left to right” list.
Scale guide changes blocking unexpectedly	Use Height / scale only instead of scale + placement.
Wrong voice	Check Subject binding, Dialogue speaker, Auto/Always state, and physical Audio inputs.
Character appears unexpectedly	Search Shot action, Timed Beats, dialogue, camera targets, frame anchors, and workflow settings for that Subject.
Camera timing warning	If Timed Beats fill the Shot and the camera moves simultaneously, use Throughout the shot or another simultaneous timing.

Quick reference

- New camera angle = New Shot
- Second action/no cut = Timed Beat
- Multiple views/same place = One Environment
- Multi-subject scale = ordered left-to-right list
- Wardrobe later = New look
- On-camera wardrobe change = action + Appearance after Shot
- Source edit = guided workflow
- Continue clip = Video Continuation

18. Privacy, Security, and Release Notes

Local-only design

The V5.14.0 release build is a standalone HTML application. The audited file contains no external JavaScript or stylesheet dependencies, no fetch/XMLHttpRequest/WebSocket/EventSource/sendBeacon network calls, no eval/new Function runtime code loading, and no external HTTP/HTTPS URLs.

Normal browser features are used for localStorage autosave, clipboard copy, user-selected JSON loading, and local JSON download.

Release summary

V5.14.0 is an audit-hardening release. It fixes invalid built-in composition enum values, corrects multi-Subject placement grammar, self-heals invalid camera enum values in migrated or auto-tidied projects, deduplicates scale mappings, expands Prompt Check coverage, counts the complete detailed_description for the word-count advisory, and removes stale dead code. The ordered multi-Subject scale/placement workflow and Basic-view connected-input controls remain intact.

Independent community tool. H3 Prompt Composer is not an official MiniMax or ComfyUI product.