


H3 Prompt Composer

V5.18.8 Full User Guide



A standalone, local-first prompt-building workspace for MiniMax H3 and ComfyUI reference workflows.

Independent community tool - not an official MiniMax or ComfyUI product.

1. Purpose and mental model

H3 Prompt Composer turns filmmaking choices into structured MiniMax H3 prompts. It does not upload media or run H3. The Pictures, Videos, and Audio clips are connected separately in ComfyUI; the Composer describes their roles and checks the resulting prompt.


A Generation is one complete H3 output. A Shot is one continuous camera view. A Timed Action Beat is a timed action window inside one continuous Shot.

Output fields

T2VA, I2VA, FL2VA, and L2VA use `integrated_multimodal_description`. Ref2VA uses six top-level sections: `subject_definitions`, `summary`, `retention_analysis`, `detailed_description`, `overall_soundscape`, and `non_diegetic_music`.

2. Quick start

- Open `H3_Prompt_Composer_V5_18_8.html` in a modern browser.
- Choose T2VA, I2VA, FL2VA, L2VA, or Ref2VA.
- For Ref2VA, set Connected ComfyUI inputs to the physical sockets available for this Generation.
- Create reusable Subjects and Environments, then build one or more Generations and Shots.
- Add camera, action, timing, dialogue, soundscape, and music only where they have a clear job.
- Resolve Prompt Check errors and review warnings before copying the prompt.



Connected input counts expose the matching Picture, Video, and Audio choices.

3. Persistent library identity vs. physical routing

V5.18.8 separates the identity of project media from its temporary H3 input socket. A library image remains 'Broken Bridge ramp detail' even when the user connects it to a different legal Picture socket for a later Generation.

A Picture number is routing for the active Environment or Generation - not a permanent global image ID.


Environment	Logical views	Temporary routing
Jungle Clearing	Wide, shed side	Pictures 5-6
Jungle Road	Wide	Picture 5
Broken Bridge	Wide, approach, ramp detail	Pictures 5-7

Cross-Environment reuse is legal because alternative Environments are not connected simultaneously. Inside one Environment, views used together must still have different Picture numbers.

Pinned and reusable naming

- Use P01_Jane_CHARACTER_SHEET.png when the file is intentionally pinned to physical Picture 1.
- Use ENV01_JungleClearing_V02_SHED_SIDE.png for a logical Environment view without a permanent slot.
- ENV03 or V03 means Environment/view numbering, not Picture 3.
- Upload order and conversational phrases such as 'picture 10' do not establish physical routing.

4. AI Setup and structured import



V5.18.8 AI Project Setup explains schema v4, media naming, and reusable Environment Picture routing.

AI Setup copies instructions and a schema-v4 starter object for an LLM. The LLM should organize the project, not write the final H3 prompt. The Composer remains responsible for subject_definitions, summary, retention_analysis, detailed_description, soundscape, music, dialogue tags, and final syntax.

Recommended AI workflow

- Give the AI the relevant project media and stable names.


- Identify explicitly pinned physical inputs; otherwise let Environment views remain slot-independent.
- Ask the AI to populate Subjects, Environments, Generations, Shots, actions, camera intent, dialogue, and ambience.
- Import the single returned JSON object through AI Setup.
- Review Environment view names and temporary Picture routing before generating.

The AI example deliberately demonstrates 5-6 / 5 / 5-7. It may reuse a slot across Environments, but it must not assign the same slot twice inside one Environment.

Migration safety

Opening a V5.18.4 project preserves its current wiring. The application does not silently rewrite user-selected ComfyUI routing. Update older projects manually when reusable routing is desired.

5. Subjects and Environments



Subject cards hold stable identities and their defining references.

Subjects may be people, creatures, vehicles, props, or other reusable visual elements. A multi-view character sheet remains one Subject. Environments hold complementary views of one stable location; target camera framing comes from the Shot, not from the source image viewpoint.

Give each reference one clear job

- Subject reference: identity and stable appearance.
- Environment view: architecture, layout, materials, decor, colors, lighting, and spatial continuity.
- Standalone Picture: composition, storyboard, wardrobe, environment state, scale/placement, lighting integration, or keyframe guidance.
- Video reference: performance, camera motion, continuation, or source footage for editing.
- Audio reference: voice timbre, exact dialogue, reused audio, music style, or sound texture.

6. Generations, Shots, and timing


Use a new Shot only for an actual cut or transition. Use Timed Action Beats for ordered events inside a continuous camera view. Every Shot should state a clear opening state before describing what changes.

Request 4-15 seconds. At 24 fps, the Composer normalizes duration to the valid 17n+5 frame sequence. A requested 10.00 seconds becomes 243 frames, or an effective 10.13 seconds.

Presence and continuity

- On-screen means visibly present.
- Off-screen means audible or nearby but intentionally not visible.
- Not present excludes the Subject.
- Auto lets the Composer infer use from action and dialogue, but explicit intent is safer when ambiguity matters.
- Continuity references should name the exact Subject, object, appearance, or environment state they preserve.


Appearance continuity



Starting look is authoritative for the opening of the active Generation.

Use a saved look for reusable wardrobe states. Use Appearance after this Shot when an on-screen change must carry forward. Earlier Generations are isolated from later edits to shared appearance history.


7. Camera Builder



Build camera intent from the start frame to the end frame, then select compatible motion.

Start with framing, primary Subject, viewpoint, height, and composition. Define an end frame only when the relationship changes. Add camera motion and timing last. Avoid stacking instructions that fight one another.


8. Scale and placement



Order every visible Subject from left to right when a visual scale/placement guide contains several figures.

Height/scale only transfers relative size, proportions, and floor-plane contact. Height/scale plus approximate placement can also transfer ordering, spacing, depth, and screen placement for a selected Shot. It does not transfer pose, wardrobe, lighting, or rough-composite artifacts.

9. Dialogue, Audio, soundscape, and music




Dialogue cards bind a speaker, delivery, language, exact words, and optional post-line action.

Bind voice references at Subject level. Speaker IDs follow actual vocal order. Off-screen speech and voiceover are distinct. Exact performed dialogue and copied audio have different retention behavior from voice-timbre references.

Sound design


- Base soundscape may continue into a new Generation.
- Additional sounds are Generation-specific.
- No scene sound means no dialogue, ambience, or physical sound.
- Non-diegetic music is audience-only and controlled separately.

10. Guided video workflows



Guided workflows generate focused editing or continuation instructions.

Insert a Subject



Preserve source timing, camera, parallax, Environment, and unaffected content while integrating the new Subject.

Background replacement and relighting

Replace or retain the background, then integrate the performer with matched directional light and reflected color. Preserve the source performance, framing, and unaffected content unless the workflow explicitly changes them.


Performance and camera transfer

Transfer only the selected acting, movement, or camera attributes while target Subject references remain authoritative for identity and appearance.

Continuation

Treat the source ending as a temporal handoff. Continue motion, camera momentum, environment, and lighting without replaying or resetting the source clip's final action.

11. Prompt Check



Errors must be fixed; warnings require intentional review; info provides context.

A clean Prompt Check means no structural conflict was detected. It does not guarantee model compliance. Check input limits, reference roles, presence, dialogue speakers, environment routing, timings, and sound claims.

12. Saving, privacy, and limits

- Auto-save uses local browser storage.
- Save .json exports an editable project; Open imports it.
- Restore previous returns to the safety backup created before destructive resets or examples.
- Copy prompt copies the active Generation; Copy all generations exports the sequence.
- The standalone HTML has no external runtime dependencies and does not send project data to a server.
- Current reference limits checked by the Composer: up to 9 images, 3 videos, 3 standalone audio clips, and 12 mixed files total: prompt ceiling 7.000 characters.

The Composer does not run MiniMax H3, upload media, configure model/sampler settings, or guarantee generation results.

Release verification

V5.18.8 retains schema v4, preserves existing wiring on migration, adds optional Environment view names, permits cross-Environment slot reuse, rejects same-Environment duplicates, and keeps all six built-in workflow prompts identical to V5.18.4.

End of guide