

H3 Prompt Composer V5.19.5

Full User Guide

Community tool documentation

August 13, 2026

1. What H3 Prompt Composer Is
2. Modes
3. Core Vocabulary
 - 3.1 Generation
 - 3.2 Shot
 - 3.3 Timed Action Beat
4. Ref2VA: The Three-Column Workflow
 - Left column - project library and inputs
 - Middle column - current Generation
 - Right column - generated prompt and review
5. Connected ComfyUI Inputs
6. Subjects and Canonical Labels
 - 6.1 Where names are still appropriate
7. Reference Jobs and Preservation Strength
 - 7.1 Preservation / retention in Basic mode
 - 7.2 Intentionally loose / flexible Picture use
8. Multi-view Environments
9. Previous-Generation Continuity
 - 9.1 Blocking continuity
 - 9.2 Multiple continuity images
10. Visual Height / Scale / Placement
 - 10.1 Height / scale only - safest
 - 10.2 Height / scale + approximate placement
11. Opening State and Main Action
 - Opening State / Blocking
 - Main Action
12. Timed Action Beats
13. Camera Builder
 - 13.1 Start Frame
 - 13.2 End Frame
 - 13.3 Camera Motion
14. Dialogue
15. Audio, Soundscape, and Music
 - Shot Sound
 - Overall soundscape
 - Non-diegetic music
16. Appearance and Wardrobe Continuity
17. Guided Video Workflows
 - 17.1 Insert a Subject into Existing Footage
 - 17.2 Background Replacement + Relighting
 - 17.3 Performance / Camera Transfer
 - 17.4 Continue an Existing Video
18. AI Project Setup
 - 18.1 Replace vs. append
19. Prompt Check
 - 19.1 Locate from Check
20. Suggestions / Prompt Coach
 - H3 adherence findings
 - Prompt Coach
 - Suggestion controls
21. H3 Adherence Field Guide
22. Phrase Bank
23. Saving, Opening, and Recovery
24. Privacy and Offline Design
25. Practical Troubleshooting
 - Prompt Check only shows one or two messages
 - Suggestions is empty
 - The model ignores a difficult instruction
 - A character refuses to turn away
 - A continuity screenshot hijacks the new composition
 - A scale diagram affects pose or blocking
 - I cannot find partially_preserved / attribute_transfer / weak_reference
26. Scope and Limitations


1. What H3 Prompt Composer Is

H3 Prompt Composer is a standalone browser tool for building structured MiniMax H3 video prompts, with particular emphasis on Ref2VA and ComfyUI reference workflows.

The application is a single HTML file. It does not install a model, upload reference media, or connect to MiniMax. It helps you express filmmaking intent as structured prompt text and catches many configuration mistakes before generation.

The central idea is simple:

- You decide the filmmaking intent.
- The Composer owns the H3 prompt skeleton.
- Prompt Check validates project structure and mappings.
- Suggestions / Prompt Coach offers non-blocking wording and relationship guidance.



V5.19.5 overview

2. Modes

The mode selector supports:

- **T2VA** - text-to-video.
- **I2VA** - generate from a required first image.
- **FL2VA** - first-frame to last-frame generation.
- **L2VA** - generate toward a required final image.
- **Ref2VA** - full-reference workflows using Subjects, Pictures, Videos, and Audio.

The largest part of the interface is devoted to Ref2VA because it requires the most explicit routing between reference inputs and the target video.

3. Core Vocabulary

3.1 Generation

A **Generation** is one complete H3 output. References are stored in the project library and can be reused across Generations.

Use **+ Same scene** when a new Generation stays in the established location. This can carry forward reusable references, appearance continuity, ambience, and music while letting you establish fresh Shot blocking.


Use **+ New location** when the scene changes location. Character identity and reusable library references remain available, but the active Environment and inherited ambience are reset so the new place can be established cleanly.

3.2 Shot

A **Shot** is one continuous camera view inside a Generation.

Use a new Shot when the camera cuts or transitions to genuinely new information. If the camera remains continuous and only the action timing changes, use Timed Action Beats instead.

V5.19.5 keeps + **Add shot** in the top toolbar and also provides + **Add Shot** below the Shot stack. The bottom control is useful in long Generations because you do not need to scroll back to the top. A new Shot is created expanded, scrolled into view, and focused on **Opening State / blocking**.



Bottom Add Shot workflow

3.3 Timed Action Beat

A **Timed Action Beat** is a timed action window inside one continuous Shot.

Timed Beats are useful when several sequential actions need clear pacing. They are not cuts. Dialogue can be linked directly to a Beat so the spoken line appears at the correct point in the Shot timeline.

4. Ref2VA: The Three-Column Workflow

The Ref2VA interface is organized around three different jobs.

Left column - project library and inputs

Store reusable Subjects and specialized references here:

- characters / people
- creatures / animals
- vehicles
- props / objects
- Environments

- Picture references
- Video references
- Audio references

The left column also defines the physical ComfyUI input counts and the active routing map.

Reference cards use conventional disclosure chevrons in V5.19.5: ► expands a collapsed card and ▼ collapses an expanded card. The + symbol is reserved for adding new things.

Middle column - current Generation

Build the target Generation one Shot at a time:

1. Scene and continuity
2. Opening state / blocking
3. Action and timing
4. Camera
5. Dialogue and on-screen text
6. Shot sound

Right column - generated prompt and review

The right column contains:

- **Prompt** - the generated H3 prompt.
- **Check** - structural and mapping validation.
- **Suggestions** - H3 adherence and Prompt Coach advice.
- **Phrases** - a static phrase bank drawn from H3-style examples.

5. Connected ComfyUI Inputs

Connected input counts represent physical H3 / ComfyUI sockets, not persistent project identities.

For example:

- <Subject 1> may be Jane for the entire project.
- Jane's character sheet may be routed through <Picture 1>.
- An Environment may use Pictures 5 and 6 in one Generation.
- A different inactive Environment may also be configured to use Picture 5 because only one of those Environment alternatives is connected at generation time.

The Composer checks the currently active file count and physical slot assignments against its encoded H3 limits.

6. Subjects and Canonical Labels

A Subject is the semantic identity that persists across the project. In Ref2VA prompt instructions, use canonical labels such as:

<Subject 1>

rather than relying on the human-readable name alone.

Preferred instructional prose:

<Subject 1> runs toward <Subject 2>.

Less explicit routing:

Jane runs toward John.

Prompt Check can recognize a known Subject name used in instructional prose and remind you of the canonical label. It can also catch malformed forms such as `Subject 1`, `<Subject1>`, `<subject 1>`, or raw `S1` when canonical reference syntax is expected.

The inline **Insert** buttons under Opening State, Main Action, Timed Beats, and Shot Sound are the safest way to insert the exact label syntax.

6.1 Where names are still appropriate

Literal dialogue and on-screen text are different. If a character actually says “Jane” or a sign contains a person’s name, that is intentional content. Prompt Check therefore avoids converting names inside spoken or displayed text.

7. Reference Jobs and Preservation Strength

Every active reference should have a clear job. A useful mental model is:

source -> contribution -> target -> boundaries

For example, a Video can supply camera motion, a Subject Picture can supply appearance, and a voice Audio reference can supply timbre. The Shot still determines the target action and framing unless one of those responsibilities has explicitly been delegated.


7.1 Preservation / retention in Basic mode

V5.19.5 makes the current preservation / transfer marker visible directly on each reference card in Basic mode.

Depending on the reference type, you may see values such as:

- `fully_preserved`
- `partially_preserved`
- `attribute_transfer`
- `weak_reference`
- `fully_copy`
- `partially_copy`
- `reference`

Editable references include **Edit in Advanced**, which switches to Advanced mode and jumps directly to the retention selector. Specialized Picture roles such as scale guidance and previous-generation continuity keep the fixed behavior that matches their intended job.



Basic-mode preservation visibility

7.2 Intentionally loose / flexible Picture use

The Composer still recommends declared roles rather than an unconstrained generic reference. Advanced users can nevertheless use a Picture more loosely when that is intentional: let the image contribute character / wardrobe / setting evidence while the current Opening State and Camera Builder establish new blocking and framing.

When doing this, make the boundaries explicit and use the Advanced preservation / transfer control rather than treating the Picture as an exact first-frame anchor.

8. Multi-view Environments

An Environment Subject can contain multiple clean views of the same location.

Use complementary views when they clarify geography or structural details, for example:

- WIDE - overall geography
- APPROACH - alternate navigation view
- SHED_SIDE - structural side view
- RAMP_DETAIL - specific construction detail

A clean Environment establishes the world. The Camera Builder establishes the target Shot viewpoint.

Previous-generation screenshots containing old characters or vehicles generally belong under Continuity References, not inside the clean Environment definition.

9. Previous-Generation Continuity

Use a previous-generation continuity Picture when a successful prior output contains state that should carry into the next Generation.

The continuity role can preserve:

- Subject starting position / blocking
- Object or vehicle position / orientation
- Environment state / local geography
- Visible appearance / physical state
- Broad scene state / spatial relationships

9.1 Blocking continuity

For blocking, think in **environment-relative spatial state**, not screenshot coordinates.

The useful information is not “Jane is 200 pixels from the left edge.” The useful information is closer to:

- Jane is beside the shed wall.
- John is one step to Jane’s right.
- Both are oriented toward the clearing.
- The Jeep remains on the road side of the location.

Unless **Preserve approximate camera framing too** is enabled, the current Shot camera reconstructs that spatial state from the new viewpoint.

9.2 Multiple continuity images

If separate screenshots establish different Subjects, the Composer can reconcile those local anchors through the clean Environment reference. This is useful when one previous-generation frame clearly establishes Jane while another clearly establishes John.

Crop or clean continuity screenshots whenever visible old state could compete with the new target.

10. Visual Height / Scale / Placement

Use text-only height relationships when possible. When a visual diagram is useful, add a **Visual height / scale / placement** Picture.

10.1 Height / scale only - safest

Scale-only mode transfers the **relative overall physical size / standing-height relationship** among the assigned Subjects.

The left-to-right Subject list is for **figure identification only**.

The scale image does not define:

- target blocking
- target spacing
- target depth
- final framing
- target pose
- target lighting

The Subjects’ own references remain authoritative for identity, anatomy, individual body proportions, wardrobe, and appearance.


10.2 Height / scale + approximate placement

This mode additionally allows the guide to influence approximate:

- left-to-right relationship
- spacing
- depth
- floor contact
- Shot blocking

The current Environment and camera remain authoritative for the actual location perspective and final composition.

Prompt Check can warn when a placement-enabled scale Picture and a previous-generation blocking Picture both influence the same Subject in the same Shot.



Scale and continuity controls

11. Opening State and Main Action

These two fields have different jobs.

Opening State / Blocking

Describe what is already true at frame zero:

- position
- orientation
- important occlusion
- who is already moving
- environment-relative starting state

Example:

<Subject 1> is already halfway across the clearing, moving toward the Jeep. <Subject 2> remains against the shed wall facing the field.

Main Action

Describe what changes after that opening state.

Example:

<Subject 2> leaves cover, runs across the clearing toward <Subject 1>, and stops beside the Jeep.

The Prompt Coach can flag Opening and Action text that repeat the same information or an extremely thin action such as <Subject 1> walks when a path or result would materially improve clarity.

12. Timed Action Beats

Use Timed Beats when one continuous Shot contains several ordered physical changes.


Prompt Check validates:

- valid start/end times
- Shot-range boundaries
- overlap between Beats
- empty Beat rows
- significant uncovered time
- dialogue linked to deleted Beats
- camera timing that is logically placed after a Beat sequence occupying the whole Shot

AI import discards completely empty imported Beats unless dialogue intentionally uses the Beat as a timing window.

13. Camera Builder

The Camera Builder is endpoint-first.



Camera Builder

13.1 Start Frame

Define:

- framing size
- primary Subject / target
- optional second or foreground Subject
- viewpoint
- camera angle
- composition
- focus detail when useful

13.2 End Frame

Keep the same frame relationship or define a different endpoint.

This is especially useful for a push, pull, zoom, orbit, pedestal, or reveal where the model needs to understand the final composition.

13.3 Camera Motion

Choose the physical move after the endpoints are defined.

The Composer validates many incompatible or stale Camera Builder states. Suggestions can also flag competing camera ownership when a Video reference is assigned to transfer camera motion while a separate target move is independently specified.

14. Dialogue

Dialogue cards separate the exact spoken words from performance direction.

The Composer owns the H3 dialogue markup. Enter the spoken words only; do not paste `<d>` wrappers into the text box.

A dialogue card can define:

- speaker
- delivery
- language
- exact line
- voice reference behavior
- off-screen / voiceover state
- carry-over across cuts
- action after the line
- optional Timed Beat relationship

Prompt Check catches many dialogue problems, including missing speakers, markup pasted into spoken words, performance directions inside the literal line, stale voice overrides, off-screen inconsistencies, and broken cross-cut dialogue pairings.

15. Audio, Soundscape, and Music

The Composer treats these as separate layers.

Shot Sound

Use for synchronized one-off physical sounds tied to the Shot, such as footsteps, impacts, door movement, engine starts, and creature roars.

Overall soundscape

Use for the continuous scene sound bed: room tone, traffic, insects, wind, sustained machinery, or other persistent ambience.

Non-diegetic music

Use for audience-only score.

When the intent is no audience-only score, choose the dedicated **No non-diegetic music** status so the Composer emits the appropriate field value rather than relying on a prose prohibition.

16. Appearance and Wardrobe Continuity

Starting look for this Generation is the authoritative opening appearance.

Use saved looks for recurring wardrobe / visible appearance states. If the appearance visibly changes inside a Shot, describe the physical transition in the Shot timeline and use the appearance-continuity controls to carry the resulting state forward.

The Composer protects earlier Generations when later appearance states are edited.

17. Guided Video Workflows

17.1 Insert a Subject into Existing Footage

Treat the source Video as the plate. Preserve timing, camera, parallax, Environment, and unaffected source content while integrating the new Subject.

Use scale/placement guidance only when apparent size or rough blocking needs extra help.

17.2 Background Replacement + Relighting

Replace the source background while preserving the performer and source framing, then rebuild environmental lighting on the performer.

Optional Lighting / Integration Pictures transfer lighting relationships without becoming pose or framing references. The workflow also supports a relight-only pass for an already-composited source.

17.3 Performance / Camera Transfer

A Video can transfer selected performance attributes such as facial expression, mouth performance, head movement, eyeline, upper-body gesture, full-body movement, and camera motion.


The target Subject reference remains authoritative for target identity / appearance.

17.4 Continue an Existing Video

Use the continuation Video as a temporal handoff. The new output starts after the source clip ends rather than re-editing the source frame for frame.

Continuity options can preserve Subject momentum, camera movement, Environment state, and prevent replaying the source ending.

18. AI Project Setup



AI Project Setup

AI Setup is an interchange workflow for external LLMs. The Composer copies concise instructions under a schema-v4 JSON structure that an LLM can fill in.

The AI can help create:

- Subjects
- Environment views
- Picture / Video / Audio references
- Generations
- Shots
- actions
- dialogue
- reference routing


After import, the Composer validates both schema structure and the generated project state.

18.1 Replace vs. append

The JSON `operation` field determines whether an import replaces the current project or appends new Generations.

For continuation work, use **Copy next-Generation context** so the external LLM receives the current state and returns append JSON rather than replacing the project.

19. Prompt Check



Prompt Check with Locate

Prompt Check is the deterministic preflight layer.

Severity meanings:

- **Error** - fix before generating.
- **Warning** - review the relationship or configuration.
- **Info** - useful context.

A clean result states that no known structural, timing, reference, dialogue, audio, camera, continuity, or mapping problems were found. The engine runs many conditional checks even when only one or two messages are visible.

19.1 Locate from Check


V5.19.5 adds **Locate** buttons to Check results when the Composer knows which field produced the problem.

Locate can jump to many Shot stages, appearance controls, reference cards, retention selectors, timing fields, sound/music status controls, and connected-input counts. The destination is expanded, scrolled into view, highlighted briefly, and focused when editable.

A project-wide warning may not have a Locate button when there is no single unique control to fix.

20. Suggestions / Prompt Coach

Suggestions never block prompt output.



Suggestions and Prompt Coach

H3 adherence findings

Examples include:

- ambiguous turn-away action
- rotation without a lens-visible end state
- weak cut / no new information
- soft audio negatives in prose
- generic quality-booster language
- detailed-description length outside the normal Ref2VA guidance range

Prompt Coach

The Prompt Coach uses field context and small deterministic vocabularies. It can flag:


- Opening and Action duplication
- very thin movement
- under-specified orientation
- thin manipulation
- ambiguous multi-Subject pronouns
- dialogue timing pressure
- competing camera instructions
- reference-role contradictions

The app is not pretending to understand arbitrary prose like an LLM. The goal is to catch common, explainable mistakes consistently.

Suggestion controls

- **Locate** - jump to the relevant field.
- **Why?** - open the related Adherence Guide section when available.
- **Dismiss** - hide advice that does not apply.

21. H3 Adherence Field Guide



H3 Adherence Field Guide

The built-in guide summarizes the prompting principles used by the Suggestions engine:

- exact skeleton, explicit descriptions
- turn-away/orientation stress cases
- soft negatives
- observable behavior vs. empty quality language
- reference jobs and ownership
- canonical labels in instructional prose
- cuts vs. camera moves
- resolution/adherence notes
- one-variable retry discipline

The guide deliberately distinguishes official guidance, converging community findings, and reproducible community tests when those evidence labels are relevant.

22. Phrase Bank

The Phrases tab is a static writing aid. It is not generated from the current prompt.

Categories include Camera, Framing, Light & grade, Physical action, Soundscape, Music, On-screen text, Continuity, and Delivery.

Clicking a phrase copies it to the clipboard as a register example. Use fragments as inspiration, not filler.

23. Saving, Opening, and Recovery

- **Auto-saved locally** - current project in browser local storage.
- **Save .json...** - opens Save As when supported; otherwise downloads normally.
- **Open** - load a saved project JSON.
- **Restore previous** - restore the safety backup created before destructive project changes.
- **Copy prompt** - copy the active Generation.
- **Copy all generations** - copy the full multi-Generation prompt sequence.

V5.19.5 migrates V5.19.4, V5.19.3, V5.19.2 TEST, V5.19.1 TEST, V5.19.0, and earlier supported project state forward.

24. Privacy and Offline Design

The audited V5.19.5 HTML is self-contained:

- no external JavaScript
- no external stylesheet
- no `fetch()`
- no `XMLHttpRequest`
- no WebSocket / EventSource
- no `sendBeacon`
- no `eval()` or downloaded executable code
- no embedded HTTP / HTTPS URLs

Normal browser capabilities are used for local storage, clipboard access, user-selected local files, native Save As where supported, and browser-download fallback.

25. Practical Troubleshooting

Prompt Check only shows one or two messages

That is normal. Check is conditional. Many validators run silently when the project state is valid. Use **Locate** when it appears instead of hunting manually for the source control.

Suggestions is empty

Also normal. Suggestions does not manufacture feedback. A blank panel means the current deterministic rules did not find a known adherence, relationship, or obvious field-clarity issue.

The model ignores a difficult instruction

Try simplifying the Shot before adding more text:

1. Confirm each active reference has one clear job.
2. Remove competing camera / blocking authorities.
3. Make the opening state explicit.

4. Describe the physical path and lens-visible end state.
5. Change one variable per retry so you can identify what helped.

A character refuses to turn away

Turn-away is a known stress case in reference-heavy generation. Consider:

- beginning a new Shot already facing away;
- describing feet, hips, shoulders, and head rotating together;
- stating the final rear-facing state explicitly;
- using a rear / three-quarter rear reference;
- separating the beat into another Generation when a Generation-level retention change would weaken identity everywhere.

A continuity screenshot hijacks the new composition

Make sure framing continuity is off when you only need world-space blocking. Crop the screenshot so old visual information that should not survive is minimized.

A scale diagram affects pose or blocking

Use **Height / scale only - safest** unless approximate placement is intentional. The left-to-right mapping identifies figures but should not control target placement in scale-only mode.

I cannot find partially_preserved / attribute_transfer / weak_reference

In Basic mode, read the **Preservation** line on the reference card and click **Edit in Advanced**. The full retention / transfer selector appears in Advanced mode. Some specialized roles use fixed markers intentionally.

26. Scope and Limitations

Prompt Check and Prompt Coach are deterministic offline systems. They cannot inspect the semantic content of your reference images or understand arbitrary prose with the flexibility of a language model.

The Composer does not guarantee model adherence. It gives H3 a cleaner, more internally consistent instruction set and helps identify common mistakes before you spend a generation.

The tool is independent community software and is not an official MiniMax or ComfyUI product.