

COMMUNITY TOOL DOCUMENTATION

H3 Prompt Composer V5.36.1

Full User Guide

Six video workflows plus a reference-guided image workflow

Guided reference routing, editing setup, and continuity tools

Visual camera planning with one active camera authority

Prompt Check and offline adherence coaching

Complete workflow reference

August 21, 2026

Standalone HTML - local browser app - no installation

Contents

- 01 What the Composer is
- 02 What changed in V5.36.1
- 03 Quick start
- 04 Modes and output grammar
- 05 Generations, Shots, and Timed Action Beats
- 06 Connected inputs and Guided Reference Setup
- 07 Subjects, labels, and reference ownership
- 08 Environments, continuity, scale, and placement
- 09 Writing a Shot
- 10 Camera Builder and Visual Planner
- 11 Dialogue, visible text, and audio
- 12 Guided video editing workflows
- 13 Targeted Edit camera reconstruction
- 14 Frame Grabber
- 15 Image mode
- 16 AI Project Setup
- 17 Prompt Check, Suggestions, and the field guide
- 18 Saving, migration, privacy, and recovery
- 19 Troubleshooting
- 20 Audit results and limitations

How to use this guide

Start with Quick Start, then use the later chapters as a field reference. The Composer owns the prompt skeleton; you supply the filmmaking intent.

1. What the Composer is

H3 Prompt Composer is a standalone browser tool for building structured MiniMax H3 prompts, with special depth for ComfyUI reference workflows. It does not run H3 or load the media connected to your workflow. It turns your filmmaking choices into prompt text and checks the project for structural conflicts before you generate.

- You decide the visible world, action, timing, camera, dialogue, references, and sound.
- The Composer owns the required field order, labels, timestamps, dialogue tags, and audio sections.
- Prompt Check handles high-confidence configuration problems; Suggestions provides non-blocking coaching.
- The complete application is one HTML file and runs locally in a modern browser.


Important boundary

The Composer describes already-connected Picture, Video, and Audio inputs. It is not a media uploader and cannot verify the contents of those files.

2. What changed in V5.36.1

V5.36.1 is the publication candidate after a broad series of workflow upgrades since the repository's previous V5.19.5 release. The release-specific refinement is a shorter, affirmative Targeted Edit camera prompt: the camera instruction leads, one optional orientation/background cue resolves the new view, and defensive mirror/world-geometry wording is removed.

- Targeted Edit now keeps exactly one camera-authoring method active: Guided controls, Visual Planner, or one custom instruction.
- The Visual Planner converts timed camera positions and path geometry into natural camera language.
- Guided Reference Setup starts from physical model inputs and assigns each one a clear job.
- Video Editing Setup now covers insertion, replacement, targeted edits, background replacement and relighting, performance/camera transfer, and continuation.
- Spatial Placement Maps, improved Environment routing, height/scale setup, and continuity tools make reference ownership more explicit.
- The local Frame Grabber creates purpose-named PNG references from video frames.
- Image mode builds a separate one-frame reference prompt and is intentionally outside the supplied video-prompt grammar.



V5.36.1 in Ref2VA mode: project inputs on the left, the active Generation in the center, and the generated prompt on the right.

3. Quick start

Step	Action
1	Open H3_Prompt_Composer_V5_36_1.html in Chrome, Edge, Firefox, or another modern browser.
2	Choose T2VA, I2VA, FL2VA, L2VA, Ref2VA, or Image mode.
3	Set the target setup and style. In Ref2VA, set Connected ComfyUI Inputs before adding references.
4	Define the Scene, Opening State, Action, Camera, Dialogue, and Shot Sound for each Shot.
5	Configure overall soundscape and audience-only music separately.
6	Review Check, then Suggestions. Fix errors; review warnings; use coaching only when it fits your intent.
7	Copy prompt and paste it into the H3 prompt field in ComfyUI.

The fastest safe Ref2VA path
Choose Ref2VA, press Guided setup, map the physical inputs, build one Shot, then review the Current Generation input map before copying the prompt.

4. Modes and output grammar

Mode	Reference behavior	Generated structure
T2VA	Text only	Three core fields; no alignment instruction
I2VA	Picture 1 is the first frame at 0.00 s	Exact first-frame instruction, then three core fields
FL2VA	Picture 1 opens; Picture 2 ends	Exact first/last alignment instruction, then three core fields
L2VA	Picture 1 is the final frame	Exact final-frame alignment instruction, then three core fields
Ref2VA	Subjects, Pictures, Videos, and Audio	Six full-reference sections
Image	One-frame reference/edit workflow	Independent image prompt; not governed by the video guides

Base video modes

```
integrated_multimodal_description: [Shot 1] ...

overall_soundscape: ...

non_diegetic_music: ...
```

I2VA, FL2VA, and L2VA place the exact alignment instruction on the first line, followed by one blank line. Shot 1 has no timestamp. Every later Shot uses a strictly increasing MM:SS.mmm cut time.

Full-reference mode

```
subject_definitions:
...

summary:
...

retention_analysis:
...

detailed_description:
...

overall_soundscape:
...

non_diegetic_music:
...
```

Ref2VA writes all six sections in English except literal dialogue, lyrics, and visible on-screen text. For ordinary generation tasks, the detailed description normally targets 350-500 English words; editing descriptions scale with the source and requested change.

5. Generations, Shots, and Timed Action Beats

Generation

A Generation is one complete H3 output. Reusable references live in the project library. Same scene carries compatible continuity forward; New location keeps reusable identities but resets location-dependent state.

Shot

A Shot is one continuous camera view. Add a Shot only for a cut or transition that introduces new information: a different subject, place, state, viewpoint, or time. A small distance or angle change is usually better as camera motion.

Timed Action Beat

A Timed Action Beat is a timed window inside one continuous Shot. Use Beats for sequential actions that need pacing without a cut. Dialogue can be linked directly to a Beat so its line appears at the right moment.

- Shot 1 never receives a timestamp.
- Later cut times must increase and remain within the effective duration.
- The requested 4-15 second duration snaps upward to H3's valid $17n+5$ frame grid at 24 fps.

6. Connected inputs and Guided Reference Setup

Connected input counts describe the physical model sockets for the active workflow, not every media file in your project. The Composer currently checks up to 9 Picture inputs, 3 Video inputs, 3 standalone Audio inputs, 12 mixed reference files total, and a 7,000-character prompt ceiling.

Guided Reference Setup images video audio Close

Plan the physical Picture, Video, and Audio inputs connected to H3. You may own many more files than the model can receive at once; this wizard defines what each connected input is reserved for. Your work stays open until you explicitly choose Close, Cancel, or Create reference setup.

1 How many model inputs are you using?
Count connected inputs, not every file in your project. H3 accepts up to nine Picture inputs.

☒ **I will use Picture inputs**
How many Picture inputs are connected?
4

☒ **I will use Video inputs**
How many Video inputs are connected?
1

☒ **I will use Audio inputs**
How many Audio inputs are connected?
1

2 What is each input reserved for?
Assign characters, environment banks, scale, continuity, composition, wardrobe, lighting, performance, or audio roles. Shot routing happens later while you build the Generation.

PICTURE INPUTS

Picture input 1 Fixed model input

Reserve this input for: Subject / character identity Subject name: Character 1 What kind of Subject is this?: Character / person

Use the same name on another Picture input only when both Images show the same Subject.

Describe Character 1
Describe permanent identity and the current temporary look separately. The optional first-shot control is only a visibility instruction.
Stable identity — what should never change

Cancel Create reference setup

Guided Reference Setup begins with the physical inputs, then assigns a semantic job to each slot.

- A Subject label is a stable semantic identity; a Picture label is a physical input slot.
- Inactive Environment alternatives may reuse Picture numbers because only the selected Environment is routed for a Generation.
- The Current Generation input map is the final load list to compare with ComfyUI before generation.
- Give every connected input one clear contribution and clear boundaries.

7. Subjects, labels, and reference ownership

Use Subject cards for visible content that persists as an identity: people, creatures, vehicles, props, and Environments. In Ref2VA instructional prose, the canonical routing identity is the exact label, such as `<Subject 1>`.

Preferred: `<Subject 1> runs toward <Subject 2>`.
Less explicit: Jane runs toward John.

Reference jobs

- Picture: frame anchor, storyboard, wardrobe, scale/placement, lighting integration, continuity, or edited keyframe.
- Video: source edit, continuation source, performance reference, camera reference, or whole-video timing structure.
- Audio: voice timbre, exact performed dialogue, copied source audio, music style, rhythm, or sound texture.

Retention markers

Visible content	Meaning
fully_preserved	The defined role is fully preserved.
partially_preserved	The reference remains active, but selected characteristics change.
attribute_transfer	Defined attributes move to another target subject.
weak_reference	Only broad style, category, composition, or atmosphere remains.

Audio relationships use fully_copy, partially_copy, reference, or weak_reference. A clean project makes each source-to-contribution-to-target relationship explicit.

8. Environments, continuity, scale, and placement

Multi-view Environments

Store complementary views of one location inside one Environment Subject. The views define a stable place and layout; the Shot still owns its target framing. Use descriptive, slot-independent filenames such as ENV03_BrokenBridge_V02_APPROACH.png unless a file is intentionally pinned to a physical Picture number.

Previous-generation continuity

Use a continuity Picture for state that should survive from a successful earlier Generation: subject blocking, object position, Environment state, appearance state, or broad spatial relationships. Crop away stale evidence that should not carry forward.

Height, scale, and placement

- Height/scale only transfers relative physical size or standing-height relationships and is the safest option.
- Height/scale plus approximate placement may also guide spacing, depth, floor contact, and blocking.
- A Spatial Placement Map is coordinates only. It does not supply finished pixels, target appearance, or lighting integration.
- Subject references remain authoritative for identity, anatomy, wardrobe, and appearance.

9. Writing a Shot

Stage	What belongs here
1 Scene and subjects	Environment, lighting, visible characters, and reference selection
2 Opening State	Where everyone is at frame zero; pose, orientation, placement, and object state
3 Action and timing	Observable motion and state changes; Timed Beats when needed
4 Camera	Framing, viewpoint, physical movement, endpoints, optics, and focus
5 Dialogue	Stable speaker, delivery, language, exact line, voice behavior, and post-line action
6 Sound	One-off synchronized impacts, footsteps, cloth, machinery, breath, or other Shot events

Describe what is visible or audible rather than writing plot summary. State the overall style and opening composition at Shot 1, then develop the action along the timeline. In keyframe modes, describe the path between the anchored frame states rather

than restating static images.


10. Camera Builder and Visual Planner

Camera Builder

The structured Camera Builder is endpoint-first: define the Start Frame, define an optional End Frame, then choose a physically compatible movement between them. Streamlined controls cover common framing and motion; Detailed controls add exact start/end geometry, composition, lens, depth of field, focus, and a custom override.

- Motion type says how the camera moves: push, pull, pan, truck, tilt, pedestal, arc, track, static, shake, POV, roll, or zoom.
- Amplitude describes the range only when it matters; medium range is normally omitted.
- Speed is added only when slow or fast pacing is meaningful.
- Keep camera wording inside natural Shot prose; avoid stacking detached labels.

Visual Planner



The Visual Planner maps timed camera positions around the subject and emits one natural camera instruction.

Place two to six timed camera positions, choose straight or arc geometry, then select subject-relative or scene-space behavior. The green arrow establishes the subject's front. Left and right remain physical subject sides when the grid follows the subject.

One camera authority

In Targeted Edit, choose Guided controls, Visual Planner, or one custom instruction. Drafts from the other methods remain dormant and do not create competing trajectories.

11. Dialogue, visible text, and audio

Speakers and dialogue

- Speaker IDs follow the order of actual vocal events: (S1), (S2), and so on. The same source keeps the same ID across Shots.
- Inside , keep only the language tag and the exact spoken words. Preserve original language and punctuation.

- Voiceover uses the exact phrase says in an off-screen voiceover and immediately states that on-screen lips remain closed.
- Use when a line crosses a cut and when speech is truncated by the end of the video.

```
<Subject 2> (S1) says: <d>[English] I get off at the next station.</d>
```

Visible text

Preserve every visible banner, sign, label, subtitle, or neon phrase exactly in English double quotation marks. Do not translate it.

Soundscape and music

- Shot Sound holds synchronized one-off sounds.
- overall_soundscape is 1-4 English sentences summarizing ambience, physical sounds, and non-verbal human sounds.
- non_diegetic_music is 1-3 English sentences describing audience-only instrumentation, tempo, rhythm, and dynamics.
- Use N/A for no non-diegetic music. Use soundscape silence only when the entire scene truly has no dialogue, ambience, or physical sound.

12. Guided video editing workflows

Workflow	Primary source relationship
Insert subject	Preserve the source plate and integrate a new subject with correct scale, grounding, occlusion, and lighting.
Replace subject	Replace an existing visible entity while inheriting source performance, timing, contact, and interactions.
Targeted edit	Change one defined element and protect unaffected content; optionally reconstruct the camera.
Background + relight	Replace the location/background and rebuild directional light and reflected color on the retained foreground.
Performance/camera transfer	Transfer selected acting, gesture, body motion, mouth performance, eyeline, dialogue, or camera behavior.
Continuation	Begin immediately after Video 1 and carry forward motion, direction, camera momentum, Environment, and lighting.

Use the Editing Setup Wizard to create only the references needed for the chosen edit. Prompt Check knows when source footage owns the camera and warns when a second camera authority is accidentally added.

13. Targeted Edit camera reconstruction

Video Editing Setup

video wizard

Close

Video 1 is the source footage. Answer only the questions that matter for this edit; the Composer will preserve source-owned camera, timing, and unaffected content automatically.

1

What kind of edit are you making?

The available choices are specific to image or video editing.

Edit type

Targeted edit of existing footage

2

Define the requested source-video change

Describe the new viewpoint, framing, or camera movement; the source timing, recorded action, performance, and environment stay authoritative.

What visible element should change? — optional when only changing the camera

the performer's jacket

Requested visible change — optional when only changing the camera

Replace only the jacket with a fitted dark-brown leather jacket.

☒ Yes — reconstruct the camera viewpoint or movement

Video 1 supplies the recorded performance and timing; your Camera instruction replaces the source camera.

How do you want to create the camera instruction?

Start with guided controls, then open the Visual Planner

Only the selected method is used in the prompt. The Composer will not combine a custom path with different dropdown settings.

Subject already in Video 1

the performer walking along the beach

Camera orientation / background — optional

The camera is on the ocean side and points inland toward the dunes and distant buildings.

Use this only when the new angle should face a particular part of the location or show a particular background.

Target framing

Wide / full shot

Camera relationship to the subject

On subject's right — show right profile

Cancel

Create editing setup


V5.36.1 Targeted Edit Setup: define the visible change, enable camera reconstruction, and choose exactly one authoring method.

Video 1 remains authoritative for the recorded performance, timing, subject motion, and location. The active camera method supplies the new shot size, physical viewpoint, camera movement, and aim. An optional orientation/background cue states what the new view should face and reveal.

Camera: use a wide right-profile view. Travel beside the subject at matching speed and keep the lens aimed at the subject. The camera faces inland; dunes, beach access, and distant buildings appear behind the subject.

The concise affirmative form is deliberate. Do not add defensive mirror language or attempt to restate the entire world's geometry. One physical relationship plus one needed background cue is clearer.

14. Frame Grabber



The local Frame Grabber captures the decoded native-resolution frame and creates a purpose-readable filename.

- Open or drop MP4, MOV, or WebM footage; the media remains on your computer.
- Step by one frame or 0.5 seconds, choose the intended Picture slot and source Generation, then select a reference purpose.
- Choose a folder for direct saving when supported, or use the browser's normal download behavior.
- For a Spatial Placement Map, capture the source frame, edit it externally, and load the rough composite as coordinates only.

15. Image mode





Image mode uses reusable reference assignments to build one coherent still-image edit or composite.

Image mode is a separate one-frame workflow. Choose a preset such as targeted edit, character sheet, storyboard, head replacement, pose/depth transfer, wardrobe/location composite, or style transformation. Assign each reference a role, state the requested change, protect everything that must remain source-authentic, and review the resulting image prompt.

Scope note

The supplied VIDEO_PROMPT_WRITING_GUIDE files govern the five video modes. Image mode is audited for runtime, layout, and internal prompt consistency, but it does not use the three-field or six-field video output grammar.


16. AI Project Setup



AI Project Setup exports instructions and schema-v4 context for an external LLM, then validates returned JSON before import.

- Copy new-project instructions for a complete first setup.
- Copy full-project context when the LLM needs the current project state.
- Copy next-Generation context to append without replacing prior Generations.
- Review physical Picture routing after import. Upload order and a filename like image 10 do not establish Picture 10 unless the slot is explicit.
- Structural schema failures block import; Prompt Check findings can be imported and fixed in the Composer.

17. Prompt Check, Suggestions, and the field guide



A clean built-in reference example: an informational frame-grid note and a passing Prompt Check.

Prompt Check

Errors are high-confidence structural failures to fix before generation. Warnings identify relationships or timing that need review. Information explains normalized state. Locate jumps to the relevant control when one clear source exists.

Suggestions and Prompt Coach

Suggestions never block output. They combine evidence-backed H3 relationship checks with deterministic field-aware coaching for duplication, weak orientation, generic quality prose, negative wording, dialogue pressure, under-specified movement, or competing reference/camera ownership.

Adherence Field Guide and Phrases

The built-in guide explains the reasoning behind adherence findings. The Phrases tab is a static register reference. It copies examples; it does not rewrite your prompt.

18. Saving, migration, privacy, and recovery

- Auto-save stores the project in local browser storage.
- Save .json opens Save As in supported Chromium browsers and falls back to a normal download elsewhere.
- Open loads a saved project JSON. Restore previous returns to the safety backup created before destructive setup, example, reset, or import actions.

- V5.36.1 migrates the V5.36.0 project and backup first, then follows the complete earlier-version fallback chain.
- The app contains no external script or stylesheet dependencies and makes no fetch, XHR, WebSocket, EventSource, sendBeacon, eval, or dynamic Function calls.

19. Troubleshooting

Symptom	What to check
Prompt Check reports missing inputs	Match Connected ComfyUI Inputs to the physical sockets you will actually connect, then review the Current Generation input map.
A frame reference hijacks composition	Check whether it should be a frame anchor, continuity Picture, attribute guide, or weak reference. Reduce its ownership.
Scale guide changes pose or blocking	Use height/scale only, or disable placement transfer. Keep Subject references authoritative for appearance.
Camera instructions conflict	Choose one camera authority. Disable camera transfer from a Video or clear Camera Builder/Visual Planner/custom wording that should be dormant.
Dialogue feels rushed	Shorten the line, enlarge its Timed Beat, adjust the cut, or use another Generation. Do not force too much speech into a short window.
Suggestions is empty	No current deterministic rule fired. This is not proof that all prose is optimal.
Model ignores a difficult turn-away	Use explicit final orientation, consider a rear/three-quarter-rear reference, or separate the stress case into its own Generation.

20. Audit results and limitations

- The exact V5.36.1 alignment instructions and three-field base-mode grammar match the supplied base guide.
- Ref2VA emits the six required sections in the correct order and follows the supplied label, retention, speaker, and audio rules.
- All six built-in workflows generate without Prompt Check errors.
- Clean browser loading produced no page exceptions, console errors, internal error bar, or external requests.
- Responsive checks at 390, 1280, 1440, 1920, and 2560 px found no document, body, or column horizontal overflow.
- The standalone HTML has no external resource tags or network/executable-code APIs listed in the privacy section.

What an audit cannot promise

A clean Prompt Check and guide-conformant prompt do not guarantee model compliance. Generation still depends on the model, reference quality, resolution, workflow settings, and the inherent difficulty of the requested motion or edit.

H3 Prompt Composer is an independent community tool and is not an official MiniMax or ComfyUI product.